
DSC 40B - Homework 09


Due: Wednesday, March 22

Write your solutions to the following problems by either typing them up or handwriting them on another piece of paper. Unless otherwise noted by the problem's instructions, show your work or provide some justification for your answer. Homeworks are due via Gradescope at 11:59 p.m.

Problem 1.

In this problem you will be asked to list the edges in the minimum spanning tree of the graph below in the order that they are added by either Prim's algorithm or Kruskal's algorithm.


In order to simplify grading, please write an edge with the *smaller* node first. For example: (u_3, u_7) instead of (u_7, u_3) . Also, when writing an edge, make sure to write the edge as a pair of nodes, not the weight of the edge. Thanks!



- Suppose Prim's algorithm is run on the above graph, using node u_1 as the starting node. List the edges of the resulting minimum spanning tree computed in the order that they are added by the algorithm.
- Suppose Kruskal's algorithm is run on the graph above. List the edges of the resulting minimum spanning tree in the order that they are added by the algorithm.

Problem 2.

The picture below shows a set of points in 2-dimensional space. A grid is provided so that you can compute the distance between points; each grid cell is 1 unit wide and 1 unit tall. You may assume that each data point is placed on a grid intersection.



Suppose a weighted distance graph G is constructed from this data set (recall that a distance graph is a complete graph whose nodes represent points in space, and whose edges are weighted by the distance between its endpoints). Then suppose that a minimum spanning tree is computed for G . What will be the weight of the largest edge in this minimum spanning tree?

Programming Problem 1.

In lecture, we saw that Kruskal's algorithm can be used to cluster a weighted graph. The name for this approach is *single linkage clustering*.

In a file named `slc.py`, write a function `slc(graph, d, k)` which accepts the following arguments:

- `graph`: An instance of `dsc40graph.UndirectedGraph`.
- `d`: A function of two arguments which takes in two nodes and returns the distance (or dissimilarity) between them.
- `k`: A positive integer describing the number of clusters which should be found.

The function should perform single linkage clustering using Kruskal's algorithm and it should return a `frozenset` of k `frozensets`, each representing a cluster of the graph.

Example:

```
>>> g = dsc40graph.UndirectedGraph()
>>> edges = [('a', 'b'), ('a', 'c'), ('c', 'd'), ('b', 'd')]
>>> for edge in edges: g.add_edge(*edge)
>>> def d(edge):
...     u, v = sorted(edge)
...     return {
...         ('a', 'b'): 1,
...         ('a', 'c'): 4,
...         ('b', 'd'): 3,
...         ('c', 'd'): 2,
...     }[(u, v)]
>>> slc(g, d, 2)
frozenset({frozenset({'a', 'b'}), frozenset({'c', 'd'})})
```


Note: to implement Kruskal's algorithm, you'll need an implementation of a Disjoint Set Forest data structure. We've uploaded a simple one here: <https://gist.github.com/elddridgejm/983d6ce03a82bf295599e9880ef02bab>

You can copy and paste this into `slc.py`, or put it in a separate file that is imported; if you do this, make sure to upload that file alongside `slc.py`.


Programming Problem 2.

Let `data` be a list of n unique real numbers. Furthermore, suppose that each number in `data` is assigned a color – it is either `'red'` or `'blue'`. Let `colors` be a list of n strings, such that `colors[i]` gives the color of `data[i]`. In all parts of this problem, you may assume for simplicity that there is at least one data point of each color.

- a) To begin, suppose that **all of the blue points are less than all of the red points**. In a file named `min_ell_theta.py`, write an efficient function called `learn_theta(data, colors)` which takes in two arguments – the lists `data` and `colors` as described above – and returns a single number θ such that all of the blue points are $\leq \theta$ and all of the red points are $> \theta$, as is depicted in the picture below. You may *not* assume that `data` is sorted. The time complexity of your algorithm should be optimal.



Now suppose that a small number of the red points are less than some blue points – that is, there is some overlap, as shown below. Assume for simplicity that the largest data point is red.



We wish to find a real number θ which “best” separates the blue points and red points. Clearly the points cannot be separated perfectly. Instead, we define a loss function $L(\theta)$ which counts the number of points which are on the wrong side of θ . More precisely:

$$L(\theta) = (\# \text{ of red points } \leq \theta) + (\# \text{ of blue points } > \theta)$$

The loss of the θ shown above is 2, since one red point is to the left of θ and one blue point is to the right. Our goal is to design an algorithm for finding a minimizer of $L(\theta)$. This is a simple instance of the machine learning task of *classification*.

- b) Also in `min_ell_theta.py`, write a function named `compute_ell(data, colors, theta)` which takes in lists `data` and `colors` as described above, as well as a floating-point number, `theta`. It should return the loss at `theta` as a floating-point number. Your algorithm should have the best possible time complexity.
- c) Also in `min_ell_theta.py`, write a function named `minimize_ell(data, colors)` which takes in `data` and `colors` and returns a floating-point number which minimizes the loss L for that particular data set. Your algorithm should have quadratic time complexity. You may assume for simplicity that the smallest data point is blue¹.
- d) Now assume that `data` is sorted (and `colors[i]` is the color of `data[i]`). In the file called `min_ell_theta.py`, write a function `minimize_ell_sorted(data, colors)` which returns a minimizer θ in linear time. Your code should satisfy the loop invariant: “After the α th iteration, `blue_gt_theta` is the number of blue points which are greater than `data[alpha - 1]`.”

For simplicity, suppose that exactly $n/2$ of the data points are `'red'`, and $n/2$ are `'blue'`.

¹Otherwise it is possible (given a special data set) for the loss to be minimized at some θ to the left of all of the data.