

# DSC 40B

## *Theoretical Foundations II*

Lecture 4 | Part 1

### **Best and Worst Cases**

## Example 1: mean

```
def mean(arr):  
    total = 0  
    for x in arr:  
        total += x  
    return total / len(arr)
```

# Time Complexity of mean

- ▶ Linear time,  $\Theta(n)$ .
- ▶ Depends **only** on the array's **size**,  $n$ , not on its actual elements.

## Example 2: Linear Search

- ▶ **Given:** an array `arr` of numbers and a target `t`.
- ▶ **Find:** the index of `t` in `arr`, or **None** if it is missing.
- ▶ **Example:** `arr = [-3, -7, 2, 9, 1, 4]`

```
def linear_search(arr, t):  
    for i, x in enumerate(arr):  
        if x == t:  
            return i  
    return None
```

## Exercise

What is the time complexity of `linear_search`?

```
def linear_search(arr, t):  
    for i, x in enumerate(arr):  
        if x == t:  
            return i  
    return None
```

# Observation

- ▶ It looks like there are *two* extreme cases...

## The Best Case

- ▶ When the target,  $t$ , is the very first element.
- ▶ The loop exits after one iteration.
- ▶  $\Theta(1)$  time?



## The **Worst** Case

- ▶ When the target,  $t$ , is not in the array at all.
- ▶ The loop exits after  $n$  iterations.
- ▶  $\Theta(n)$  time?

# Time Complexity

- ▶ `linear_search` can take vastly different amounts of time on two inputs of the **same size**.
  - ▶ Depends on **actual elements** as well as size.
- ▶ It has no single, overall time complexity.
- ▶ Instead we'll report **best** and **worst** case time complexities.

# Best Case Time Complexity

- ▶ How does the time taken in the **best case** grow as the input gets larger?

## Definition

Define  $T_{\text{best}}(n)$  to be the **least** time taken by the algorithm on any input of size  $n$ .

The asymptotic growth of  $T_{\text{best}}(n)$  is the algorithm's **best case asymptotic time complexity**.

## Best Case

- ▶ In `linear_search`'s **best case**,  $T_{\text{best}}(n) = c$ , no matter how large the array is.
- ▶ The **best case time complexity** is  $\Theta(1)$ .

# Worst Case Time Complexity

- ▶ How does the time taken in the **worst case** grow as the input gets larger?

## Definition

Define  $T_{\text{worst}}(n)$  to be the **most** time taken by the algorithm on any input of size  $n$ .

The asymptotic growth of  $T_{\text{worst}}(n)$  is the algorithm's **worst case asymptotic time complexity**.

## Worst Case

- ▶ In the worst case, `linear_search` iterates through the entire array.
- ▶ The **worst case time complexity** is  $\Theta(n)$ .



## Exercise

What are the best case and worst case time complexities of the following code?

```
def foo(arr):  
    n = len(arr)  
    for x in arr:  
        for y in arr:  
            if x + y == 10:  
                return sum(arr)
```

## Best Case

- ▶ When the first element is 5, so  $x + y == 10$ .
- ▶ `sum(arr)` takes  $\Theta(n)$  time.
- ▶ Exits, taking  $\Theta(n)$  time in total.

# Worst Case

- ▶ No two elements sum to 10.
- ▶ Has to loop over all  $\Theta(n^2)$  pairs.
- ▶ Worst case time complexity:  $\Theta(n^2)$ .
- ▶ **Note:** it's not  $\Theta(n^3)$ , since the `sum(arr)` only runs once!

## Caution!

- ▶ The best case is never: “the input is of size one”.
- ▶ The best case is about the **structure** of the input, not its **size**.
- ▶ Not always constant time! Example: sorting.

## Note

- ▶ An algorithm like `linear_search` doesn't have **one single** time complexity.
- ▶ An algorithm like `mean` does, since the best and worst case time complexities coincide.

## Main Idea

Reporting **best** and **worst** case time complexities gives us a richer of the performance of the algorithm.

# DSC 40B

## *Theoretical Foundations II*

Lecture 4 | Part 2

**Average Case**

# Time Taken, Typically

- ▶ Best case and worst case can be **misleading**.
  - ▶ Depend on a **single good/bad input**.
- ▶ How much time is taken, typically?
- ▶ **Idea:** compute the average time taken over all possible inputs.



# Recall: The Expectation

- The expected value of a random variable  $X$  is:

$$\sum_x x \cdot P(X = x)$$

winnings	probability
\$ 0	50%
\$ 1	30%
\$ 10	18%
\$ 50	2%

Expected winnings:

# Average Case

- ▶ We'll compute the expected time over all cases:

$$T_{\text{avg}}(n) = \sum_{\text{case} \in \text{all cases}} P(\text{case}) \cdot T(\text{case})$$

- ▶ Called the **average case time complexity**.

# Strategy for Finding Average Case

- ▶ **Step 0:** Make assumption about distribution of inputs.
- ▶ **Step 1:** Determine the possible cases.
- ▶ **Step 2:** Determine the probability of each case.
- ▶ **Step 3:** Determine the time taken for each case.
- ▶ **Step 4:** Compute the expected time (average).

# Example: Linear Search

- Recall **linear search**:

```
def linear_search(arr, t):  
    for i, x in enumerate(arr):  
        if x == t:  
            return i  
    return None
```

- Best case? Worst case?

## Example: Linear Search

- ▶ What is the **average case time complexity** of **linear search**?

## **Step 0: Assume input distribution**

- ▶ We must assume something about the input.
- ▶ Example: Target must be in array, equally-likely to be any element, no duplicates.
- ▶ This is usually given to you.

# Step 1: Determine the Cases

- Example: linear search.

Case 1: target is first element

Case 2: target is second element

⋮

Case  $n$ : target is  $n$ th element

Case  $n + 1$ : target is not in array

## Step 2: Case Probabilities

- ▶ What is the probability that we see each case?
  - ▶ Example: what is the probability that the target is the  $k$ th element?
- ▶ This is where we use assumptions from Step 0.



# Example

- ▶ **Assume:** target is in the array exactly once, equally-likely to be any element.
- ▶ Each case has probability  $1/n$ .

## Step 3: Case Times

- ▶ Determine time taken in each case.
- ▶ Example: linear search.
  - ▶ Let's say it takes time  $c$  per iteration.

Case 1: time  $c$

Case 2: time  $2c$

$\vdots$

Case  $i$ : time  $c \cdot i$

$\vdots$

Case  $n$ : time  $c \cdot n$

## Step 4: Compute Expectation

$$T_{\text{avg}}(n) = \sum_{i=1}^n P(\text{case } i) \cdot T(\text{case } i)$$

# Average Case Time Complexity

- ▶ The **average case** time complexity<sup>1</sup> of **linear search** is  $\Theta(n)$ .

---

<sup>1</sup>Under these assumptions on the input!

# Note

- ▶ **Worst case** time complexity is still useful.
- ▶ Easier to calculate.
- ▶ Often same as average case (but not always!)
- ▶ Sometimes worst case is very important.
  - ▶ Real time applications, time complexity attacks

## Note

- ▶ **Hard** to make realistic assumptions on input distribution.
- ▶ Example: linear search.
  - ▶ Is it realistic to assume  $t$  is in array?
  - ▶ If not, what is the probability that it *is* in the array?

## Exercise

Suppose we change our assumptions:

- ▶ The target has a 50% chance of being in the array.
- ▶ If it is in the array, it is equally-likely to be any element.

What is the average case complexity now?

# DSC 40B

## *Theoretical Foundations II*

Lecture 4 | Part 3

**Average Case in Movie Problem**



## Recall: The Movie Problem

- ▶ **Given:** an array `movies` of movie durations, and the flight duration `t`
- ▶ **Find:** two movies whose durations add to `t`.
  - ▶ If no two movies sum to `t`, return **None**.

```
def find_movies(movies, t):  
    n = len(movies)  
    for i in range(n):  
        for j in range(i + 1, n):  
            if movies[i] + movies[j] == t:  
                return (i, j)  
    return None
```

# Time Complexity

- ▶ Best case:  $\Theta(1)$ 
  - ▶ When the first pair of movies checked equals target.
- ▶ Worst case:  $\Theta(n^2)$ 
  - ▶ When no pair of movies equals target.

## “Average” Case?

- ▶ The best and worst cases are **extremes**.
- ▶ How much time is taken, *typically*?
  - ▶ That is, when the target pair is not the first checked nor the last, but somewhere in the middle.

## Exercise

How much time do you expect `find_movies` to take on a typical input?

- ▶  $\Theta(1)$
- ▶  $\Theta(n^2)$
- ▶ Something in between, like  $\Theta(n)$

# The Movie Problem

```
def find_movies(movies, t):  
    n = len(movies)  
    for i in range(n):  
        for j in range(i + 1, n):  
            if movies[i] + movies[j] == t:  
                return (i, j)  
    return None
```

# Time Complexity

- ▶ Best case:  $\Theta(1)$
- ▶ Worst case:  $\Theta(n^2)$
- ▶ Average case:  $\Theta(?)$

## Step 0: Assume input distribution

- ▶ Suppose we are told that:
  - ▶ There is a unique pair of movies that add to  $t$ .
  - ▶ All pairs are equally likely.



## Step 1: Determine the Cases

- ▶ Case  $\alpha$ : the  $\alpha$ th pair checked sums to  $t$ .
- ▶ Each pair of movies is a case.
- ▶ There are  $\binom{n}{2}$  cases.

## Step 2: Case Probabilities

- ▶ **Assume:** there is a *unique* pair that adds to  $t$ .
- ▶ **Assume:** all pairs are equally likely.
- ▶ Probability of any case:  $\frac{1}{\binom{n}{2}} = \frac{2}{n(n-1)}$

## Step 3: Case Time

- ▶ How much time is taken for a particular case?
- ▶ Example, suppose the movies  $a$  and  $b$  sum to the target.
- ▶ How long does it take to find this pair?

```
1 def find_movies(movies, t):  
2     n = len(movies)  
3     for i in range(n):  
4         for j in range(i + 1, n):  
5             if movies[i] + movies[j] == t:  
6                 return (i, j)  
7     return None
```

### Exercise

Roughly much time is taken (how many times does line 5 run) if the  $\alpha$ th pair checked sums to the target?

## **Step 4: Compute Expectation**

## Average Case

- ▶ The average case time complexity of `find_movies` is  $\Theta(n^2)$ .
- ▶ Same as the worst case!

## Note

- ▶ We've seen two algorithms where the average case = the worst case.
- ▶ Not always the case!
- ▶ Interpretation: the worst case is not too extreme.

# DSC 40B

## *Theoretical Foundations II*

Lecture 4 | Part 4

**Expected Time Complexity**



# Example: Contrived Algorithm

```
def wibble(n):  
    # generate random number between 0 and n  
    x = np.random.randint(0, n)  
  
    if x == 0:  
        for i in range(n):  
            print('Unlucky!')  
    else:  
        print('Lucky!')
```

## Exercise

How much time does wibble take *on average*?

# Random Algorithms

- ▶ This algorithm is *randomized*.
- ▶ The time it takes is also *random*.
- ▶ What is the **expected time**?

# Average Case vs. Expected Time

- ▶ With average case complexity, a probability distribution on inputs is specified.
- ▶ Now, the randomness is *in the algorithm itself*.
- ▶ Otherwise, the analysis is very similar.

## Step 1: Determine the cases

```
def wibble(n):  
    x = np.random.randint(0, n)  
  
    if x == 0:  
        for i in range(n):  
            print('Unlucky!')  
    else:  
        print('Lucky!')
```

► Case 1:  $x == 0$

► Case 2:  $x \neq 0$

## Step 2: Determine case probabilities

```
def wibble(n):  
    x = np.random.randint(0, n)  
  
    if x == 0:  
        for i in range(n):  
            print('Unlucky!')  
    else:  
        print('Lucky!')
```

►  $P(\text{Case 1}) = 1/n$

►  $P(\text{Case 2}) = (n - 1)/n$

## Step 3: Determine case times

```
def wibble(n):  
    x = np.random.randint(0, n)  
  
    if x == 0:  
        for i in range(n):  
            print('Unlucky!')  
    else:  
        print('Lucky!')
```

► Case 1:  $\Theta(n)$

► Case 2:  $\Theta(1)$

## **Step 4: Compute expectation**

- ▶ Compute expected time:

# Expected Time

- ▶ This was a contrived example.
- ▶ Some important algorithms involve randomness!
  - ▶ Quicksort
  - ▶ We'll see alg. for median with  $\Theta(n)$  expected time.



# DSC 40B

## *Theoretical Foundations II*

Lecture 4 | Part 5

### **Lower Bound Theory**

# Imagine...

- ▶ You write a simple algorithm to solve a problem.
- ▶ You analyze time complexity and find it is  $\Theta(n^2)$ .
- ▶ You ask yourself: *can I do better than  $\Theta(n^2)$ ?*
- ▶ Or: *What is the best time complexity possible?*

# Doing Better

- ▶ How can you know what you don't know?
- ▶ You can argue that *any* algorithm for solving the problem *must* take at least a certain amount of time in the worst case.

# Example: Minimum

- ▶ Problem: Find minimum in array of length  $n$ .
- ▶ *Any* algorithm has to check all  $n$  numbers in the worst case.
  - ▶ Or else the number not checked could have been the smallest!
- ▶ Takes at least linear ( $\Omega(n)$ ) time.
  - ▶ **No algorithm** for the min can have worst case of  $<$  linear time.

## Definition

A **theoretical lower bound** is a lower bound on the worst-case time complexity of **any algorithm** solving a particular problem.

## Main Idea

No algorithm's worst case can be better than theoretical lower bound.

# Loose Lower Bounds

- ▶  $\Omega(\log n)$ ,  $\Theta(\sqrt{n})$  and  $\Theta(1)$  are also theoretical lower bounds for finding the minimum.
- ▶ But no algorithm can exist which has a worst case of  $\Theta(\log n)$ ,  $\Theta(\sqrt{n})$ , or  $\Theta(1)$ .
- ▶ This bound is **loose**. Not super useful.

# Tight Lower Bounds

- ▶ A lower bound is **tight** if there exists an algorithm with that worst case time complexity.
- ▶ That algorithm is (in a sense) **optimal**.



# How to find a TLB

- ▶ Argument from completeness:
  - ▶ The algorithm might not be correct if it doesn't check  $k$  things, so the time is  $\Omega(k)$ .
- ▶ Argument from I/O:
  - ▶ If the output is an array of size  $k$ , time taken is  $\Omega(k)$
- ▶ More sophisticated arguments...

# **Tight Bounds can be difficult to find**

- ▶ Often require sophisticated combinatorial arguments outside of the scope of DSC 40B.

# Assumptions make problems easier

- ▶ The TLB for finding a minimum changes if we assume that the array is sorted.

## Exercise

Consider these two problems:

1. Find the min of a sorted array.
2. Given a target  $t$  and a sorted array, determine whether  $t$  is in the array.

Find tight theoretical lower bounds for each problem.

## Main Idea

When coming up with an algorithm, first try to find a tight TLB. Then try to make an algorithm which has that worst-case complexity. If you can, it's **optimal!**

# DSC 40B

## *Theoretical Foundations II*

Lecture 4 | Part 6

### **Case Study: Matrix Multiplication**

# It's Important

- ▶ Matrix multiplication is a *very* common operation in machine learning algorithms.
- ▶ **Estimate:** 75% - 95% of time training a neural network is spent in matrix multiplication.

# Recall

- ▶ If  $A$  is  $m \times p$  and  $B$  is  $p \times n$ , then  $AB$  is  $m \times n$ .
- ▶ The  $ij$  entry of  $AB$  is

$$(AB)_{ij} = \sum_{k=1}^p a_{ik} b_{kj}$$



# Recall

$$(AB)_{ij} = \sum_{k=1}^p a_{ik} b_{kj}$$

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{pmatrix} \begin{pmatrix} 5 & -1 \\ 1 & 7 \\ -2 & -3 \end{pmatrix} = \begin{pmatrix} & & \\ & & \\ & & \end{pmatrix}$$

# Naïve Algorithm

- ▶ This algorithm is relatively straightforward to code up.

```
def mmul(A, B):  
    """  
    A is (m x p) and B is (p x n)  
    """  
  
    m, p = A.shape  
    n = B.shape[1]  
  
    C = np.zeros((m, n))  
  
    for i in range(m):  
        for j in range(n):  
            for k in range(p):  
                C[i,j] += A[i,k] * B[k, j]  
  
    return C
```

# Time Complexity

- ▶ The naïve algorithm takes time  $\Theta(mnp)$ .
- ▶ If both matrices are  $n \times n$ , then  $\Theta(n^3)$  time.
- ▶ **Cubic!**

# Cubic Time Complexity

- ▶ The largest problem size that can be solved, if a basic operation takes 1 nanosecond.

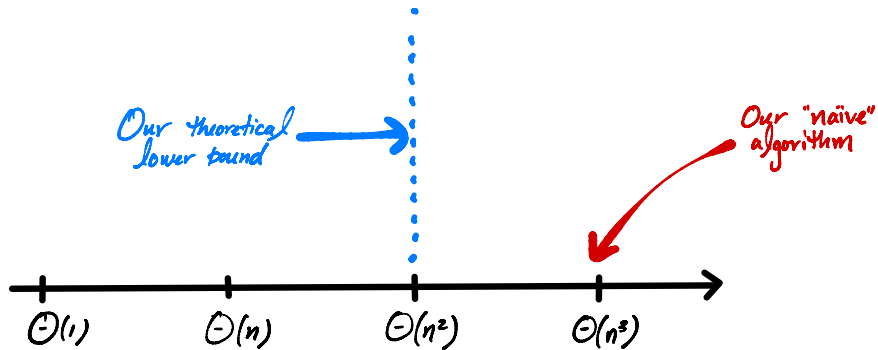
1 s	10 m	1 hr
1,000	6,694	15,326

# The Question

- ▶ Can we do better?
- ▶ How fast can we possibly multiply matrices?

# Theoretical Lower Bound


- ▶ If  $A$  and  $B$  are  $n \times n$ ,  $C$  will have  $n^2$  entries.
- ▶ Each entry must be filled:  $\Omega(n^2)$  time.
- ▶ That is, matrix multiplication must take at least quadratic time.
- ▶ Is this bound **tight**? Can it be increased?






# Strassen's Algorithm

- ▶ Cubic was as good as it got...
- ▶ ...until Strassen, 1969.
- ▶ Time complexity:  $\Theta(n^{\log_2 7}) = \Theta(n^{2.8073})$



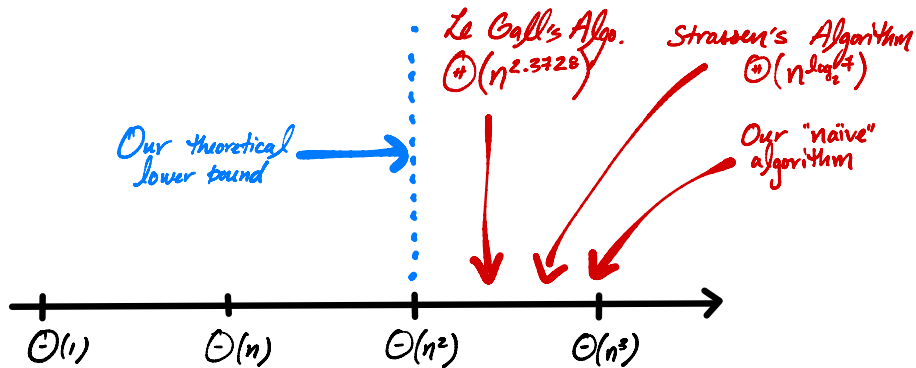


# Currently

- ▶ The fastest<sup>2</sup> known matrix multiplication algorithm is due to Le Gall.
- ▶  $\Theta(n^{2.3728639})$  time.

---

<sup>2</sup>In terms of asymptotic time complexity.



## Interestingly...

- ▶ No one knows what the lowest possible time complexity is.
- ▶ It could be  $\Theta(n^2)$ !
- ▶ The “best” matrix multiplication algorithm is probably still undiscovered.

# Irony

- ▶ There are many matrix multiplication algorithms.
- ▶ How fast is numpy's matrix multiply?

# Irony

- ▶ There are many matrix multiplication algorithms.
- ▶ How fast is numpy's matrix multiply?
- ▶  $\Theta(n^3)$ .



# Why?

- ▶ Strassen *et al.* have better asymptotic complexity.
- ▶ But much (much!) larger “hidden constants”.
- ▶ Remember, which is better for small  $n$ :  $999,999n^2$  or  $n^3$ ?

# Optimization

- ▶ Numpy, most others use **highly optimized** cubic time algorithms<sup>3</sup>

---

<sup>3</sup>The constant  $c$  in  $T(n) = cn^3 + \dots$  is actually much less than 1, as can be verified by timing.

## Main Idea

No one knows what the lowest possible time complexity of matrix multiplication is, and some algorithms are approaching  $\Theta(n^2)$ .

But most useful implementations take  $\Theta(n^3)$  time.