


DSC 40B

Theoretical Foundations II

Lecture 16 | Part 1

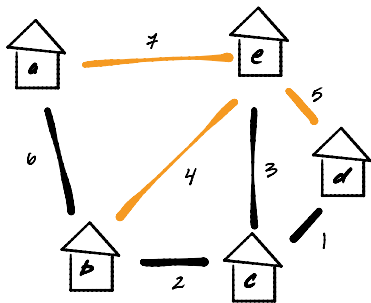
Minimum Spanning Trees

Today's Problem



- ▶ Choose a set of dirt roads to pave so that:
 - ▶ can get between any two buildings only on paved roads;
 - ▶ total cost is minimized.
- ▶ Solution: compute a **minimum spanning tree**.

Today's Problem



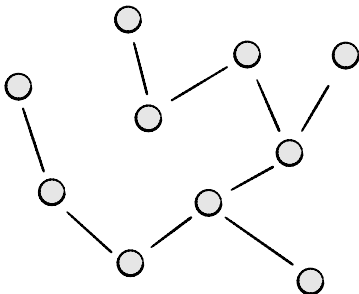
- ▶ Choose a set of dirt roads to pave so that:
 - ▶ can get between any two buildings only on paved roads;
 - ▶ total cost is minimized.
- ▶ Solution: compute a **minimum spanning tree**.

Trees

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ it is acyclic.

Example: a **tree**.

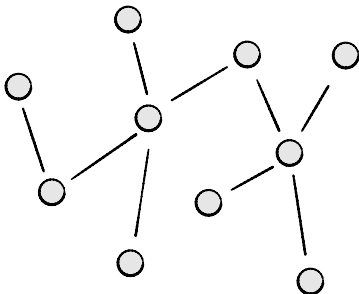


Trees

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ it is acyclic.

Example: a **tree**.

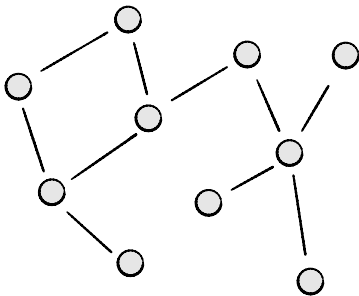


Trees

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ it is acyclic.

Example: **not** a tree.

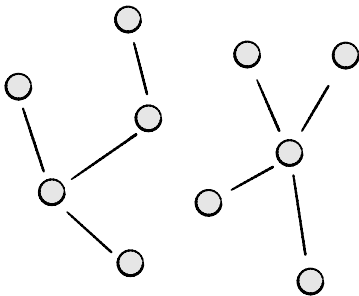


Trees

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ it is acyclic.

Example: **not** a tree.

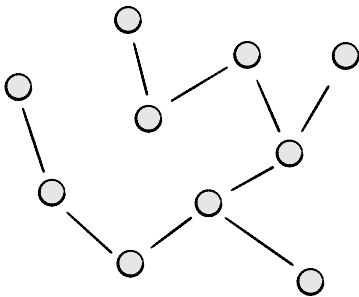


Trees: Equivalent Definition

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ $|E| = |V| - 1$.

Example: a **tree**.




Trees: Equivalent Definition

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ $|E| = |V| - 1$.

Example: a **tree**.




Trees: Equivalent Definition

An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ $|E| = |V| - 1$.

Example: **not** a tree.




Trees: Equivalent Definition

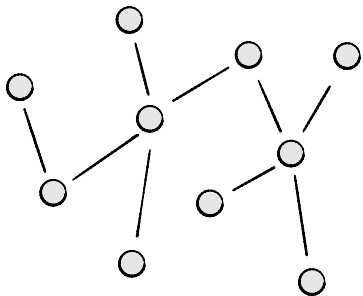
An undirected graph $T = (V, E)$ is a **tree** if

- ▶ it is connected; and
- ▶ $|E| = |V| - 1$.

Example: **not** a tree.



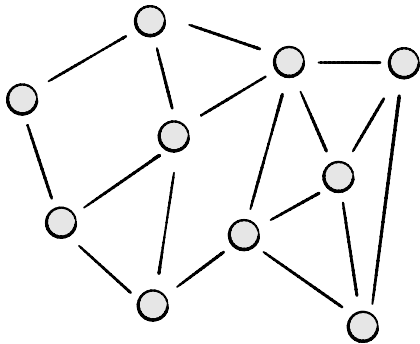
Tree Properties



- ▶ There is a unique simple path between any two nodes in a tree.
- ▶ Adding a new edge to a tree creates a cycle (no longer a tree).
- ▶ Removing an edge from a tree disconnects it (no longer a tree).

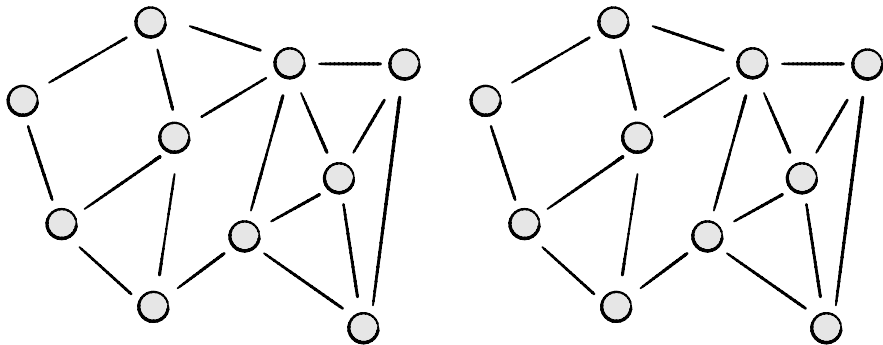
Spanning Trees

Let $G = (V, E)$ be a **connected** graph. A **spanning tree** of G is a tree $T = (V, E_T)$ with the same nodes as G , and a subset of G 's edges.




Many Spanning Trees

The same graph can have many spanning trees.



Spanning Tree Cost

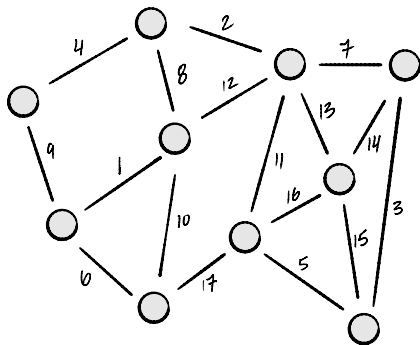
If $G = (V, E, \omega)$ is a weighted undirected graph, the **cost** (or **weight**) of a spanning tree is the total weight of the edges in the spanning tree.



Cost:

Spanning Tree Cost

Different spanning trees of the same graph can have different costs.



Cost:


Minimum Spanning Tree

- ▶ The **minimum spanning tree** problem is as follows:
 - ▶ GIVEN: A weighted, undirected graph $G = (V, E, \omega)$.
 - ▶ COMPUTE: a spanning tree of G with minimum cost (i.e., minimum total edge weight).
- ▶ For a given graph, the MST may not be unique.

Exercise

Suppose the edges of a graph $G = (V, e, \omega)$ all have the same weight. How can we compute an MST of the graph?

Today's Problem



- ▶ Choose a set of dirt roads to pave so that:
 - ▶ can get between any two buildings only on paved roads;
 - ▶ total cost is minimized.
- ▶ Solution: compute a **minimum spanning tree**.

MSTs in Data Science?

- ▶ Do we need to find MSTs in data science?
- ▶ Actually, yes! (Next lecture)

DSC 40B

Theoretical Foundations II

Lecture 16 | Part 2

Prim's Algorithm

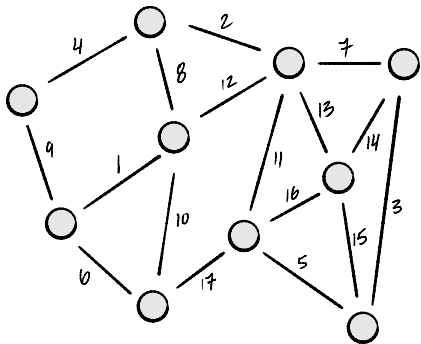
Building MSTs

- ▶ How do we build a MST efficiently?
- ▶ We'll adopt a **greedy** approach.
 - ▶ Build a tree edge-by-edge.
 - ▶ At every step, doing what looks best at the moment.
- ▶ This strategy isn't guaranteed to work in all of life's situations, but it works for building MSTs.

Two Greedy Approaches

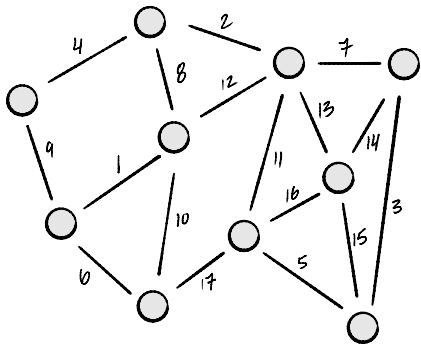
- ▶ We'll look at two greedy algorithms:
 - ▶ Today: Prim's Algorithm
 - ▶ Next time: Kruskal's Algorithm
- ▶ Differ in the order in which edges are added to tree.
- ▶ Also differ in time complexity.

Prim's Algorithm, Informally



- ▶ Start by picking any node to add to “tree”, T .
- ▶ While T is not a spanning tree, greedily add **lightest** edge from a node in T to a node not in T .
 - ▶ “lightest” = edge of smallest weight

Prim's Algorithm, Informally

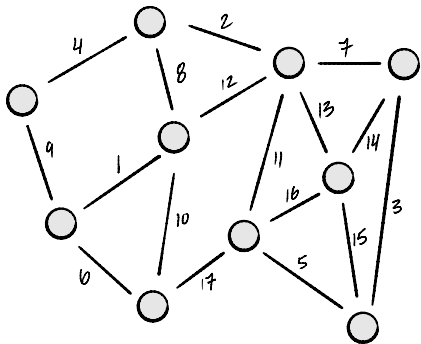


- ▶ Start by picking any node to add to “tree”, T .
- ▶ While T is not a spanning tree, greedily add **lightest** edge from a node in T to a node not in T .
 - ▶ “lightest” = edge of smallest weight
- ▶ **Is this guaranteed to work?** Yes, as we’ll see.

Prim's Algorithm, Equivalently

- ▶ For each node u , store:
 - ▶ estimated cost of adding node to tree;
 - ▶ estimated “predecessor” v in the tree.
- ▶ At each step,
 - ▶ Find node with smallest estimated cost.
 - ▶ Add to tree T by including edge with estimated “predecessor”.
 - ▶ Update cost of neighbors.
- ▶ Same as adding lightest edge from T to outside T at every step!

Prim's Algorithm, Equivalently



- ▶ While T is not a tree:
 - ▶ find the node $u \notin T$ with smallest cost
 - ▶ add the edge between u and its estimated “predecessor” to T
 - ▶ update estimated cost/pred. of u ’s neighbors which aren’t already in tree.

Recall: Priority Queues

- ▶ How do we efficiently find node with smallest cost?
- ▶ Priority Queues:
 - ▶ `PriorityQueue(priorities)`: creates priority queue from dictionary whose values are priorities.
 - ▶ `.extract_min()`: removes and returns key with smallest value.
 - ▶ `.decrease_priority(key, value)`: changes key's value.
- ▶ We'll use a priority queue to hold nodes not yet added to tree.

```
def prim(graph, weight):  
    tree = UndirectedGraph()  
  
    estimated_predecessor = {node: None for node in graph.nodes}  
    cost = {node: float('inf') for node in graph.nodes}  
    priority_queue = PriorityQueue(cost)  
  
    while priority_queue:  
        u = priority_queue.extract_min()  
        if estimated_predecessor[u] is not None:  
            tree.add_edge(estimated_predecessor[u], u)  
        for v in graph.neighbors(u):  
            if weight(u, v) < cost[v] and v not in tree.nodes:  
                priority_queue.decrease_priority(v, weight(u, v))  
                cost[v] = weight(u, v)  
                estimated_predecessor[v] = u  
    return tree
```

Prim and Dijkstra

- ▶ This is a lot like Dijkstra's Algorithm for s.p.d.!
- ▶ Both: at each step, extract node with smallest cost, update its edges. (Prim: only those edges to nodes not in tree).
- ▶ Dijkstra update of (u, v) :

$$\text{cost}[v] = \min(\text{cost}[v], \text{cost}[u] + \text{weight}(u, v))$$

- ▶ Prim update of (u, v) :

$$\text{cost}[v] = \min(\text{cost}[v], \text{weight}(u, v))$$

DSC 40B

Theoretical Foundations II

Lecture 16 | Part 3

Time Complexity

Time Complexity

- ▶ A priority queue can be implemented using a **heap**.
- ▶ If a **binary min-heap** is used:
 - ▶ `PriorityQueue(est)` takes $\Theta(V)$ time.
 - ▶ `.extract_min()` takes $O(\log V)$ time.
 - ▶ `.decrease_priority()` takes $O(\log V)$ time.

Time Complexity

```
def prim(graph, weight):  
    tree = UndirectedGraph()  
  
    estimated_predecessor = {node: None for node in graph.nodes}  
    cost = {node: float('inf') for node in graph.nodes}  
    priority_queue = PriorityQueue(cost)  
  
    while priority_queue:  
        u = priority_queue.extract_min()  
        if estimated_predecessor[u] is not None:  
            tree.add_edge(estimated_predecessor[u], u)  
        for v in graph.neighbors(u):  
            if weight(u, v) < cost[v] and v not in tree.nodes:  
                priority_queue.decrease_priority(v, weight(u, v))  
                cost[v] = weight(u, v)  
                estimated_predecessor[v] = u  
  
    return tree
```

Time Complexity

- ▶ Using a **binary heap**...
- ▶ Overall: $\Theta(V \log V + E \log V)$.
- ▶ Since graph is assumed connected, $E = \Omega(V)$.
- ▶ So this simplifies to $\Theta(E \log V)$.

Fibonacci Heaps

- ▶ A priority queue can be implemented using a **heap**.
- ▶ If a **Fibonacci min-heap** is used:
 - ▶ `PriorityQueue(est)` takes $\Theta(V)$ time.
 - ▶ `.extract_min()` takes $\Theta(\log V)$ time¹.
 - ▶ `.decrease_priority()` takes $O(1)$ time.

¹Amortized.

Time Complexity

```
def prim(graph, weight):  
    tree = UndirectedGraph()  
  
    estimated_predecessor = {node: None for node in graph.nodes}  
    cost = {node: float('inf') for node in graph.nodes}  
    priority_queue = PriorityQueue(cost)  
  
    while priority_queue:  
        u = priority_queue.extract_min()  
        if estimated_predecessor[u] is not None:  
            tree.add_edge(estimated_predecessor[u], u)  
        for v in graph.neighbors(u):  
            if weight(u, v) < cost[v] and v not in tree.nodes:  
                priority_queue.decrease_priority(v, weight(u, v))  
                cost[v] = weight(u, v)  
                estimated_predecessor[v] = u  
  
    return tree
```

Time Complexity

- ▶ Using a **Fibonacci heap**...
- ▶ Overall: $\Theta(V \log V + E)$.

Fibonacci vs. Binary Heaps

- ▶ Using Fibonacci heaps improves time complexity when graph is dense.
- ▶ E.g., if $E = \Theta(V^2)$:
 - ▶ Prim's with Fibonacci: $\Theta(E) = \Theta(V^2)$
 - ▶ Prim's with binary: $\Theta(E \log E) = \Theta(V^2 \log V)$.
- ▶ But Fibonacci heaps are **hard** to implement; have large constants.
- ▶ Binary heaps used more in practice despite complexity.

DSC 40B

Theoretical Foundations II

Lecture 16 | Part 4

Correctness of Prim's Algorithm

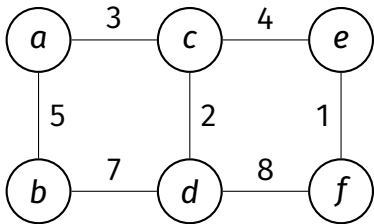
Being Greedy

- ▶ At every step, we add the lightest edge.
- ▶ Is this “safe”?

Being Greedy

- ▶ At every step, we add the lightest edge.
- ▶ Is this “safe”?
- ▶ Yes! This is guaranteed to find an MST.

Promising Subtrees



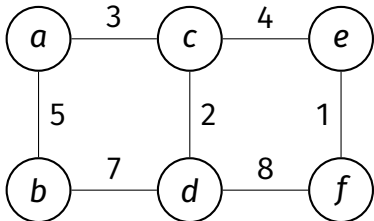
- ▶ Let $G = (V, E, \omega)$ be a weighted graph.
- ▶ A subgraph $T' = (V', E')$ is **promising** if it is “part” of some MST.
 - ▶ That is, it is an “MST in progress”
 - ▶ Not necessarily a tree!
- ▶ That is, there exists an MST $T = (V, E_{\text{mst}})$ such that $E' \subset E_{\text{mst}}$.
- ▶ Hint: a “promising subtree” where $V' = V$ is an MST!

Main Idea

Prim's starts with a promising subtree T . At each step, adds lightest edge from a node within T to a node outside of T .

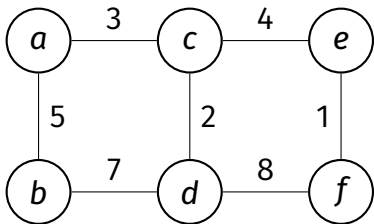
We'll show each new edge results in a larger promising subtree. Eventually the promising subtree becomes a full MST.

Claim



- ▶ Let $G = (V, E, \omega)$ be a weighted graph.
- ▶ Suppose $T' = (V', E')$ is a promising subtree for an MST of G .
- ▶ Let $e = (u, v)$ be a **lightest edge** from a node in T' to a node outside of T' . (Prim).
- ▶ Then adding (u, v) to T' results in another **promising subtree**.

Proof



- ▶ Suppose T_{mst} is an MST that includes T' .
- ▶ If T_{mst} includes e , we're done: $T' + e$ is promising.
- ▶ If it doesn't include e , it must have an edge f that connects T' to rest of the graph.
- ▶ Swap f with e in T_{mst} . The result is a tree, and it must be a MST since $\omega(e) \leq \omega(f)$.
- ▶ So there is an MST that contains $T' + e$.