

DSC 190

DATA STRUCTURES & ALGORITHMS

Lecture 17 | Part 1

Today's Lecture

Massive Sets

- ▶ You've collected 1 billion tweets.¹
- ▶ **Goal:** given the text of a new tweet, is it already in the data set?

¹This is about two days of activity.

Membership Queries

- ▶ We want to perform a **membership query** on a collection of strings.
- ▶ Hash tables support $\Theta(1)$ membership queries.
- ▶ **Idea:** so let's use a hash table (Python: `set`).

Problem: Memory

- ▶ How much memory would a **set** of 1 billion strings require?
- ▶ Assume average string has 100 ASCII characters.
 $(8 \text{ bits per char}) \times (100 \text{ chars}) \times 1 \text{ billion} = 100 \text{ gigabytes}$
- ▶ That's way too large to fit in memory!

Today's Lecture

- ▶ **Goal:** fast membership queries on massive data sets.
- ▶ Today's answer: **Bloom filters**.

DSC 190

DATA STRUCTURES & ALGORITHMS


Lecture 17 | Part 2

Bit Arrays

The Challenge

- ▶ We want to perform membership queries on a massive collection (too large to fit in memory).
- ▶ We want to remember which elements are in the collection...
- ▶ ...*without* actually storing all of the elements.
- ▶ From hash tables to Bloom filters in 3 steps.

First Stop: Hash Tables



s	hash(s)
"surf"	3
"sand"	8
"data"	5
"sun"	1
"beach"	5

Memory Usage

- ▶ **Problem:** we're storing all of the elements.
- ▶ Why? To resolve collisions.
- ▶ **Fix:** ignore collisions.

Second Stop: Hashing Into Bit Arrays

0	1	2	3	4	5	6	7	8	9
0	1	0	1	0	1	0	0	1	0

s	hash(s)
"surf"	3
"sand"	8
"data"	5
"sun"	1
"beach"	5

- Use a bit array `arr` of size `c`.
- **Insertion:** Set `arr[hash(x)] = 1`.
- **Query:** Check if `arr[hash(x)] = 1`.

Second Stop: Hashing Into Bit Arrays

0	1	2	3	4	5	6	7	8	9
0	1	0	1	0	1	0	0	1	0

s	hash(s)
"surf"	3
"sand"	8
"data"	5
"sun"	1
"beach"	5

- Use a bit array `arr` of size `c`.
- **Insertion:** Set `arr[hash(x)] = 1`.
- **Query:** Check if `arr[hash(x)] = 1`.
- Can be **wrong**!

False Positives

0	1	2	3	4	5	6	7	8	9
0	1	0	1	0	1	0	0	1	0

s	hash(s)
"surf"	3
"sand"	8
"data"	5
"sun"	1
"beach"	5

- ▶ Query can return **false positives**.
- ▶ e.g.,
`hash("ucsd") == 3`
- ▶ Cannot return false negatives.

Memory Usage

- ▶ Requires c bits, where c is size of the bit array.
- ▶ False positive rate depends on c .
 - ▶ c is small \rightarrow more collisions \rightarrow more errors
 - ▶ c is large \rightarrow fewer collisions \rightarrow fewer errors
- ▶ **Tradeoff:** get more accuracy at cost of memory.

False Positive Rate

- ▶ What is the probability of a false positive?
- ▶ Suppose there are c buckets, and we've inserted n elements so far.
- ▶ We query an object x that we haven't seen before.
- ▶ False positive $\Leftrightarrow \text{arr}[\text{hash}(x)] == 1$.

False Positive Rate

- ▶ Assume **hash** assigns bucket uniformly at random.
 - ▶ If $x \neq y$ then, $\mathbb{P}(\text{hash}(x) = \text{hash}(y)) = 1/c$
- ▶ Prob. that first element does not collide with x : $1 - 1/c$.
- ▶ Prob. that first two do not collide: $(1 - 1/c)^2$.
- ▶ Prob. that all n elements do not collide: $(1 - 1/c)^n$.

False Positive Rate

- ▶ Hint: for large z , $(1 - 1/z)^z \approx \frac{1}{e}$

- ▶ So the probability of no collision is:

$$(1 - 1/c)^n = [(1 - 1/c)^c]^{n/c} \approx e^{-n/c}$$

- ▶ This is the probability of no false positive.
- ▶ Probability of false positive upon querying x :
 $\approx 1 - e^{-n/c}$

False Positive Rate

- ▶ For fixed query, probability of false positive:
 $\approx 1 - e^{-n/c}$.
 - ▶ n : number of elements stored
 - ▶ c : size of array (number of bits)
- ▶ Randomness is over choice of hash function.
 - ▶ Once hash function is fixed, the result is always the same.

Fixing False Positive Rate

- ▶ Suppose we'll tolerate false positive rate of ε .
- ▶ Assume that we'll store around n elements.
- ▶ We can choose c :

$$1 - e^{-n/c} = \varepsilon \quad \implies \quad c = -\frac{n}{\ln(1 - \varepsilon)}$$

Example

- ▶ Suppose we want $\leq 1\%$ error.
- ▶ Previous slide says our bit array needs to be 100 times larger than number of elements stored.²
- ▶ Memory when $n = 10^9$: 1 billion bits $\times 100 = 12.5$ GB.
- ▶ Can we do better?

²We could have guessed this, huh?

DSC 190

DATA STRUCTURES & ALGORITHMS

Lecture 17 | Part 3

Bloom Filters

Wasted Space

- ▶ Suppose we want $\leq 1\%$ error.
- ▶ Our bit array needs to be 100 times larger than number of elements stored.
- ▶ That's a lot of **wasted space**!

Third Stop: Multiple Hashing

- ▶ **Idea:** use several smaller bit arrays, each with own hash function.

Third Stop: Multiple Hashing

0	1	2	3	4	5	6	7	8	9
0	1	0	1	0	1	0	0	1	0
0	1	2	3	4	5	6	7	8	9
0	0	0	0	1	0	1	1	0	1

s	hash ₁ (s)	hash ₂ (s)
"surf"	3	7
"sand"	8	7
"data"	5	4
"sun"	1	9
"beach"	5	6

- Use k bit arrays of size c , each with own independent hash function.
- **Insertion:** Set
 $\text{arr}_1[\text{hash}_1(x)] = 1,$
 $\text{arr}_2[\text{hash}_2(x)] = 1,$
...,
 $\text{arr}_k[\text{hash}_k(x)] = 1.$

Third Stop: Multiple Hashing

0	1	2	3	4	5	6	7	8	9
0	1	0	1	0	1	0	0	1	0
0	1	2	3	4	5	6	7	8	9
0	0	0	0	1	0	1	1	0	1

s	hash_1(s)	hash_2(s)
"surf"	3	7
"sand"	8	7
"data"	5	4
"sun"	1	9
"beach"	5	6

- Use k bit arrays of size c , each with own independent hash function.
- **Query:** Return **True** if **all** of
 $\text{arr}_1[\text{hash}_1(x)] = 1$,
 $\text{arr}_2[\text{hash}_2(x)] = 1$,
...,
 $\text{arr}_k[\text{hash}_k(x)] = 1$.
- Example:
 $\text{hash}_1(\text{"hello"}) == 3$,
 $\text{hash}_2(\text{"hello"}) == 2$

Intuition

- ▶ False positive occurs only if false positive in **all** tables.
- ▶ This is pretty unlikely.
- ▶ If false positive rate in one table is small (but not tiny), probability false positive in all tables is still tiny.

More Formally

- ▶ Probability of false positive in first table:
 $\approx 1 - e^{-n/c}$.
- ▶ Probability of false positive in all k tables:
 $\approx (1 - e^{-n/c})^k$.
- ▶ Example: if $c = 4n$ and $k = 3$, error rate is $\approx 1\%$.
- ▶ Uses only $12 \times n$ bits, as opposed to $100 \times n$ from before.

Last Stop: Bloom Filters

- ▶ How many different bit arrays do we use? (What is k ?)
- ▶ How large should they be? (What is c ?)
- ▶ **Bloom filters**: use k hash functions, but only one medium-sized array.

Last Stop: Bloom Filters

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
0	1	0	1	0	1	1	0	1	1	0	0	0	1	0	1	0	1	0	0	0

s	hash ₁ (s)	hash ₂ (s)
"surf"	13	17
"sand"	8	6
"data"	15	1
"sun"	1	3
"beach"	5	9

► Use one bit arrays of size c , but k hash functions.

► **Insertion:** Set
 $\text{arr}[\text{hash}_1(x)] = 1$,
 $\text{arr}[\text{hash}_2(x)] = 1$,
...,
 $\text{arr}[\text{hash}_k(x)] = 1$.

Last Stop: Bloom Filters

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
0	1	0	1	0	1	1	0	1	1	0	0	0	1	0	1	0	1	0	0	0

s	hash_1(s)	hash_2(s)
"surf"	13	17
"sand"	8	6
"data"	15	1
"sun"	1	3
"beach"	5	9

- Use one bit arrays of size c , but k hash functions.
- **Query:** Return **True** if **all** of
 $\text{arr}[\text{hash}_1(x)] = 1$,
 $\text{arr}[\text{hash}_2(x)] = 1$,
...,
 $\text{arr}[\text{hash}_k(x)] = 1$.
- Example:
 $\text{hash}_1(\text{"hello"}) == 3$,
 $\text{hash}_2(\text{"hello"}) == 2$

Intuition

- ▶ Multiple hashing allows bit arrays to be smaller.
- ▶ Even more efficient: let them share memory.
- ▶ “Overlaps” are just collisions; we can handle them.

False Positive Rate

- ▶ Consider querying new, unseen object x .
- ▶ We'll look at k bits.
 - ▶ $\text{arr}[\text{hash}_1(x)], \dots, \text{arr}[\text{hash}_k(x)]$.
- ▶ Fix one bit. What is the chance that it is already one?

False Positive Rate

- ▶ Probability of bit being zero after first element inserted: $(1 - 1/c)^k$
- ▶ After second element inserted: $(1 - 1/c)^{2k}$
- ▶ After all n elements inserted: $(1 - 1/c)^{nk}$
- ▶ And:

$$(1 - 1/c)^{nk} = [(1 - 1/c)^c]^{nk/c} \approx e^{-nk/c}$$

False Positive Rate

- ▶ Probability of bit being **one** after n elements inserted:

$$1 - e^{-nk/c}$$

- ▶ For a false positive, all k bits (for each hash function) need to be one.
- ▶ Assuming independence,³ probability of false positive:

$$(1 - e^{-nk/c})^k$$

³Only true approximately. If this bit was set, some other bit was not.

Minimizing False Positives

- For a fixed n and c , the number of hash functions k which minimizes the false positive rate is

$$k = \frac{c}{n} \ln 2$$

- Plugging this into the error rate:

$$\varepsilon = (1 - e^{-nk/c})^k \quad \implies \quad \ln \varepsilon = -\frac{c}{n} (\ln 2)^2$$

- If we fix ε , then $c = -n \ln \varepsilon / (\ln 2)^2$

Summary: Designing Bloom Filters

- ▶ Suppose we wish to store n elements with ε false positive rate.
- ▶ Allocate a bit array with $c = -n \ln \varepsilon / (\ln 2)^2$ bits.
- ▶ Pick $k = \frac{c}{n} \ln 2$ hash functions.

Example

- ▶ Let $n = 10^9$, $\varepsilon = 0.01$.
- ▶ We need $c \approx 9.5n \rightarrow 10n$ bits = 1.25 GB.
- ▶ We choose $k = \frac{9.5n}{n} \ln 2 \rightarrow 7$ hash functions.

DSC 190

DATA STRUCTURES & ALGORITHMS

Lecture 17 | Part 4

Bloom Filters in Practice

Applications

- ▶ A cool data structure.
- ▶ Most useful when data is huge or memory is small.

Application #1

- ▶ De-duplicate 1 billion strings, each about 100 bytes.
- ▶ Memory required for `set`: 100 gigabytes.
- ▶ Instead:
 - ▶ Loop through data, reading one string at a time.
 - ▶ If not in Bloom filter, write it to file.
- ▶ With 1% error rate, takes 1.25 GB.

Application #2

- ▶ A ***k*-mer** is a substring of length k in a DNA sequence:

"GATTACATATAGGTGTCGA"

- ▶ Useful: does a long string have a given k -mer?
- ▶ There are a *massive* number of possible k -mers.
 - ▶ 4^k , to be precise.
 - ▶ Example: there are over 10^{18} 30-mers.
- ▶ Slide window of size k over sequence, store each substring in Bloom filter.

Application #2

- ▶ Human genome is a 725 Megabyte string, 2.9 billion characters.
- ▶ To store all k -mers, each character stored k times.
- ▶ Storing 30-mers in **set** would take $30 \times 725 \text{ MB} \approx 22 \text{ GB}$.
- ▶ By “forgetting” the actual strings, Bloom filter (1% false positive) takes

2.9 billion bits \approx 360 megabytes

Application #3

- ▶ Suppose you have a massive database on disk.
- ▶ Querying the database will take a while, since it has to go to disk.
- ▶ Build a Bloom filter, keep in memory.
 - ▶ If Bloom filter says x not in database, don't perform query.
 - ▶ Otherwise, perform DB query.
- ▶ Speeds up time of “misses”.

Limits

- ▶ Bloom filters are useful in certain circumstances.
- ▶ But they have disadvantages:
 - ▶ Need good idea of size, n , ahead of time.
 - ▶ There are false positives.
 - ▶ The elements are not stored (can't iterate over them).
- ▶ Often a **set** does just fine, with some care.

Example

- ▶ Suppose you have 1 billion tweets.
- ▶ Want to de-duplicate them by **tweet ID** (64 bit number).
- ▶ Total size: 8 gigabytes.
- ▶ I have 4 GB RAM. Should I use a Bloom filter?

De-duplication Strategy

- ▶ Design a hash function that maps each tweet ID to $\{1, \dots, 8\}$.
- ▶ Loop through tweet IDs one-at-a-time, hash, write to file:
`hash(x) == 3` → write to `data_3.txt`
- ▶ Read in each file, one-at-a-time, de-duplicate with `set`, write to `output.txt`