

# DSC 190

## Machine Learning: Representations

Lecture 14 | Part 1

### Basic Backpropagation

# Computing the Gradient

- ▶ To train a neural network, we can use gradient descent.
- ▶ Involves computing the gradient of the cost function.
- ▶ **Backpropagation** is one method for efficiently computing the gradient.

# The Gradient

$$\begin{aligned}\nabla_{\vec{w}} C(\vec{w}) &= \nabla_{\vec{w}} \frac{1}{n} \sum_{i=1}^n (f(\vec{x}^{(i)}; \vec{w}) - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n \nabla_{\vec{w}} (f(\vec{x}^{(i)}; \vec{w}) - y_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n 2(f(\vec{x}^{(i)}; \vec{w}) - y_i) \nabla_{\vec{w}} f(\vec{x}^{(i)}; \vec{w})\end{aligned}$$

# Interpreting the Gradient

$$\nabla_{\vec{w}} C(\vec{w}) = \frac{1}{n} \sum_{i=1}^n 2 \left( f(\vec{x}^{(i)}; \vec{w}) - y_i \right) \nabla_{\vec{w}} f(\vec{x}^{(i)}; \vec{w})$$

- ▶ The gradient has one term for each training example,  $(\vec{x}^{(i)}, y_i)$
- ▶ If prediction for  $\vec{x}^{(i)}$  is good, contribution to gradient is small.
- ▶  $\nabla_{\vec{w}} f(\vec{x}^{(i)}; \vec{w})$  captures how sensitive  $f(\vec{x}^{(i)})$  is to value of each parameter.

# The Chain Rule

- ▶ Recall the **chain rule** from calculus.

- ▶ Let  $f, g : \mathbb{R} \rightarrow \mathbb{R}$

- ▶ Then:

$$\frac{d}{dx} f(g(x)) = f'(g(x)) \cdot g'(x)$$

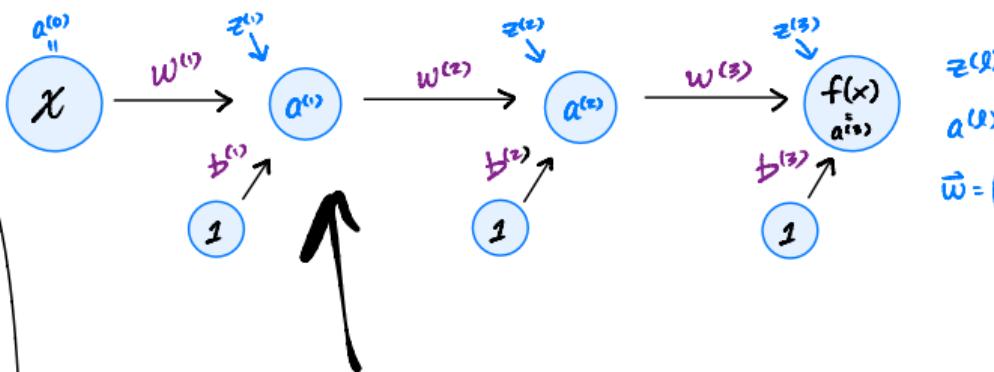
- ▶ Alternative notation:  $\frac{d}{dx} f(g(x)) = \frac{df}{dg} \frac{dg}{dx}(x)$

$$a^{(2)} = \sigma(w^{(2)}a^{(1)} + b^{(2)})$$

# The Chain Rule for NNs

$$f(x) = \sigma(w^{(3)}a^{(2)} + b^{(3)})$$

$$\begin{aligned} Df &= \\ &= \left( \frac{\partial f}{\partial w^{(1)}}, \frac{\partial f}{\partial w^{(2)}}, \dots, \frac{\partial f}{\partial b^{(3)}} \right) \end{aligned}$$



$$z^{(l)} = w^{(l)}a^{(l-1)} + b^{(l)}$$

$$a^{(l)} = \sigma(z^{(l)})$$

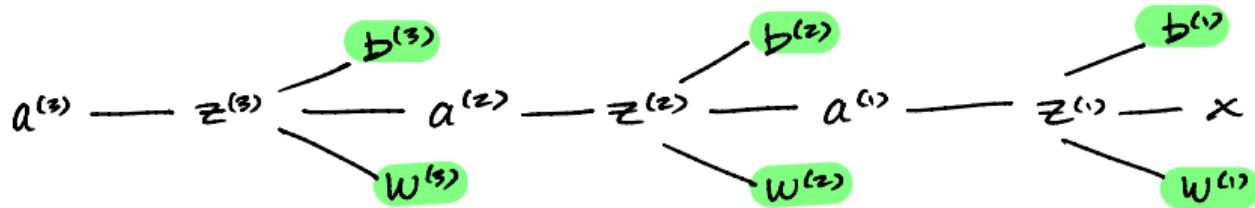
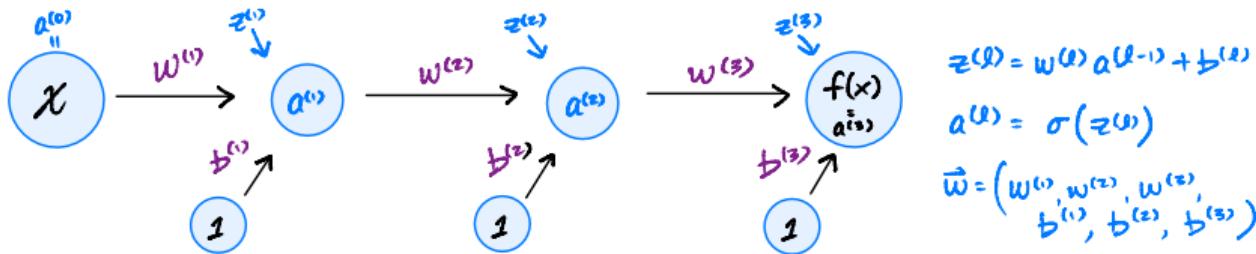
$$\bar{w} = \begin{pmatrix} w^{(1)}, w^{(2)}, w^{(3)}, \\ b^{(1)}, b^{(2)}, b^{(3)} \end{pmatrix}$$

$$z^{(l)} = w^{(l)}x + b^{(l)} \quad \text{"raw output"}$$

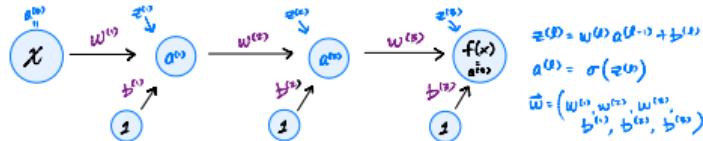
$\sigma$  "activation fn"

$$a^{(l)} = \sigma(z^{(l)}) \quad \text{"actual output"}$$

# Computation Graphs

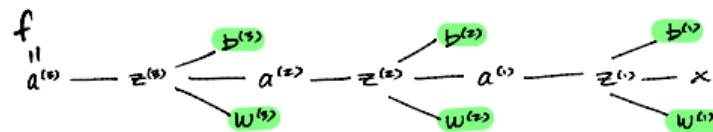


# Example



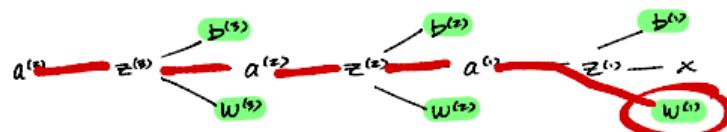
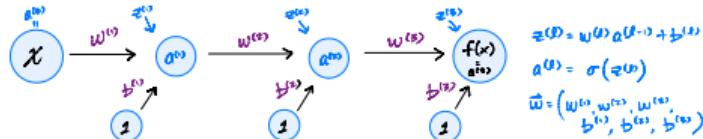
$$\begin{aligned}z^{(3)} &= \theta^{(3)} = w^{(3)} a^{(2)} + b^{(3)} \\a^{(3)} &= \sigma(z^{(3)}) \\w &= (w^{(1)}, w^{(2)}, w^{(3)}, b^{(1)}, b^{(2)}, b^{(3)})\end{aligned}$$

$$\frac{\partial f}{\partial w^{(3)}} = \frac{\partial a^{(3)}}{\partial w^{(3)}} = \underbrace{\frac{\partial a^{(3)}}{\partial z^{(3)}}}_{\sigma'(z^{(3)})} \underbrace{\frac{\partial z^{(3)}}{\partial w^{(3)}}}_{a^{(2)}}$$



$$\begin{aligned}\frac{\partial f}{\partial w^{(3)}} &= \frac{\partial a^{(3)}}{\partial w^{(3)}} & a^{(3)} &= \sigma(z^{(3)}) \\&= \frac{\partial}{\partial w^{(3)}} \sigma(z^{(3)}) & z^{(3)} &= w^{(3)} a^{(2)} + b^{(3)} \\&= \sigma'(z^{(3)}) \frac{\partial}{\partial w^{(3)}} z^{(3)} \\&= \sigma'(z^{(3)}) \frac{\partial}{\partial w^{(3)}} [w^{(3)} a^{(2)} + b^{(3)}] \\&= \sigma'(z^{(3)}) a^{(2)}\end{aligned}$$

# Example



$$\frac{\partial a^{(3)}}{\partial w^{(1)}} = \underbrace{\frac{\partial a^{(3)}}{\partial z^{(3)}}}_{\sigma'(z^{(3)})} \underbrace{\frac{\partial z^{(3)}}{\partial a^{(2)}}}_{w^{(3)}} \underbrace{\frac{\partial a^{(2)}}{\partial z^{(2)}}}_{\sigma'(z^{(2)})} \underbrace{\frac{\partial z^{(2)}}{\partial a^{(1)}}}_{w^{(2)}} \underbrace{\frac{\partial a^{(1)}}{\partial z^{(1)}}}_{\sigma'(z^{(1)})} \underbrace{\frac{\partial z^{(1)}}{\partial w^{(1)}}}_{w^{(1)}} \quad z^{(l)} = w^{(l)} x + b^{(l)}$$

$$\frac{\partial a^{(3)}}{\partial w^{(2)}} = \frac{\partial a^{(3)}}{\partial z^{(3)}} \frac{\partial z^{(3)}}{\partial a^{(2)}} \frac{\partial a^{(2)}}{\partial z^{(2)}} \frac{\partial z^{(2)}}{\partial w^{(2)}}$$

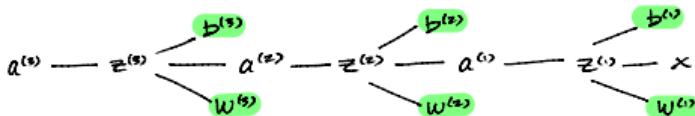
# General Formulas

- ▶ Derivatives are defined recursively
- ▶ Easy to compute derivatives for early layers if we have derivatives for later layers.
- ▶ This is **backpropagation**.

$$\frac{\partial f}{\partial w^{(3)}} = \frac{\partial f}{\partial a^{(3)}} \cdot \frac{\partial a^{(3)}}{\partial z^{(3)}} \cdot \frac{\partial z^{(3)}}{\partial w^{(3)}}$$

$$\frac{\partial f}{\partial w^{(l)}} = \frac{\partial f}{\partial a^{(l)}} \cdot \frac{\partial a^{(l)}}{\partial z^{(l)}} \cdot \frac{\partial z^{(l)}}{\partial w^{(l)}}$$

$$\frac{\partial f}{\partial a^{(l)}} = \frac{\partial f}{\partial a^{(l+1)}} \cdot \frac{\partial a^{(l+1)}}{\partial z^{(l+1)}} \cdot \frac{\partial z^{(l+1)}}{\partial a^{(l)}}$$

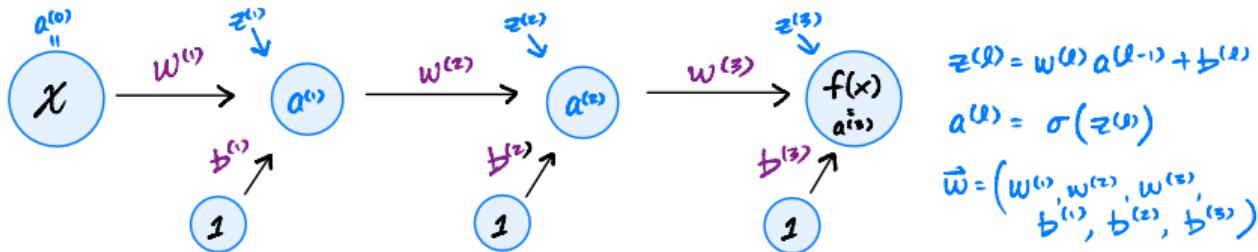


# Warning

- ▶ The derivatives depend on the network **architecture**
  - ▶ Number of hidden nodes / layers
- ▶ Backprop is done automatically by your NN library

# Backpropagation

Compute the derivatives for the last layers first; use them to compute derivatives for earlier layers.



# DSC 190

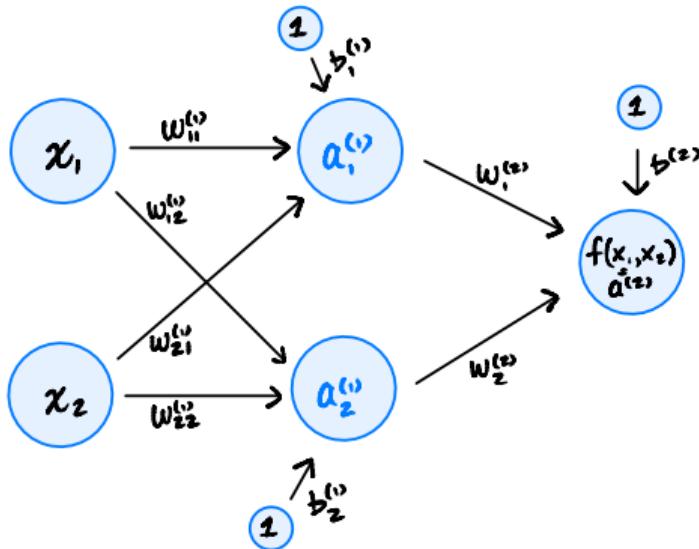
## *Machine Learning: Representations*

Lecture 14 | Part 2

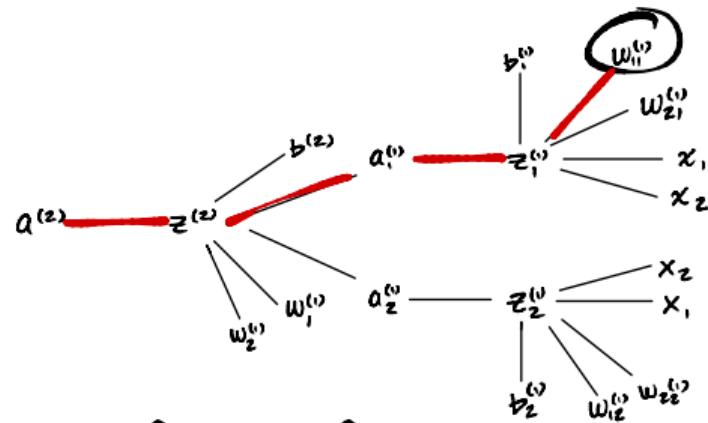
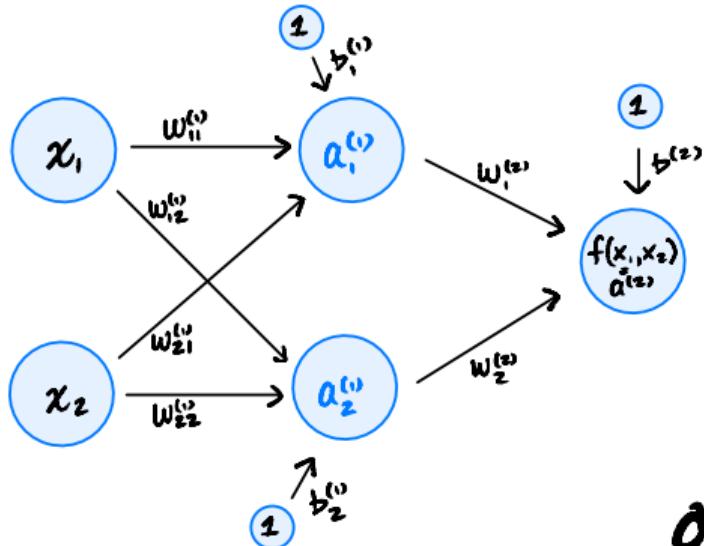
### A More Complex Example

# Complexity

- The strategy doesn't change much when each layer has more nodes.



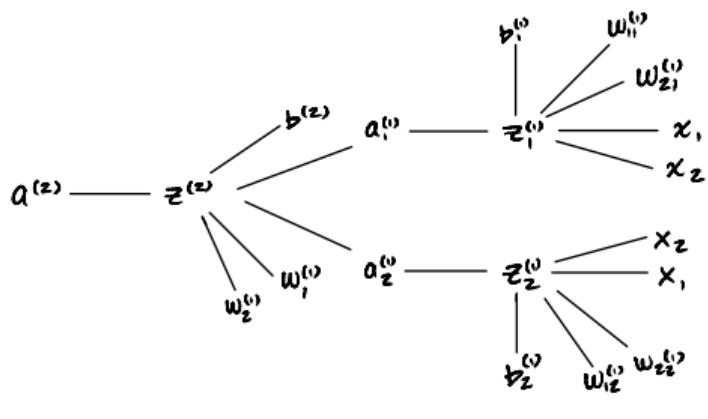
# Computational Graph



$$\frac{\partial a^{(2)}}{\partial w_{11}^{(1)}} = \frac{\partial a^{(2)}}{\partial z^{(2)}} \frac{\partial z^{(2)}}{\partial a_1^{(1)}} \frac{\partial a_1^{(1)}}{\partial z_1^{(1)}} \frac{\partial z_1^{(1)}}{\partial w_{11}^{(1)}}$$

# Example

# General Formulas



$$\frac{\partial f}{\partial w_{ij}^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot \frac{\partial a^{(\ell)}}{\partial z^{(\ell)}} \cdot \frac{\partial z^{(\ell)}}{\partial w_{ij}^{(\ell)}}$$

$$\frac{\partial f}{\partial a^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell+1)}} \cdot \frac{\partial a^{(\ell+1)}}{\partial z^{(\ell+1)}} \cdot \frac{\partial z^{(\ell+1)}}{\partial a^{(\ell)}}$$

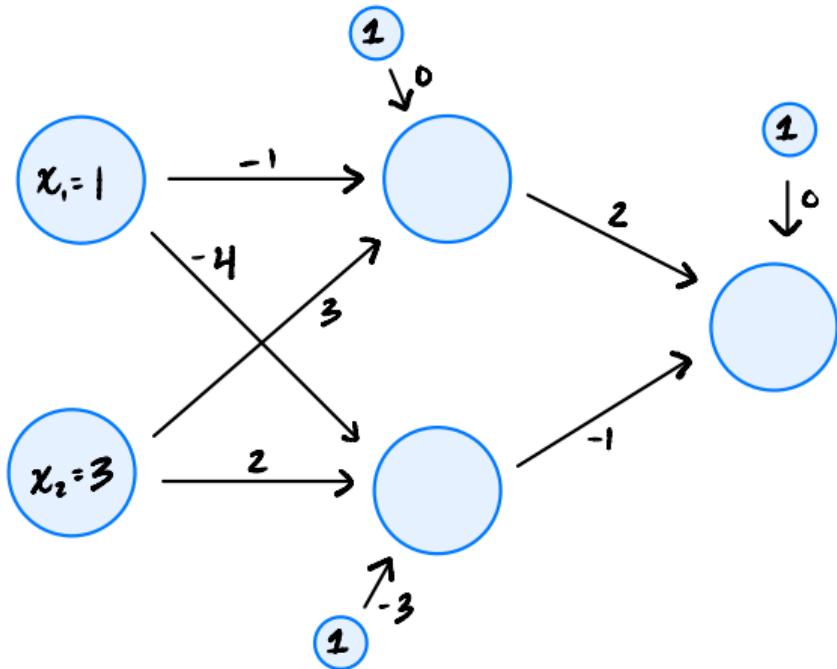
# DSC 190

## Machine Learning: Representations

Lecture 14 | Part 3

### Intuition Behind Backprop

# Intuition



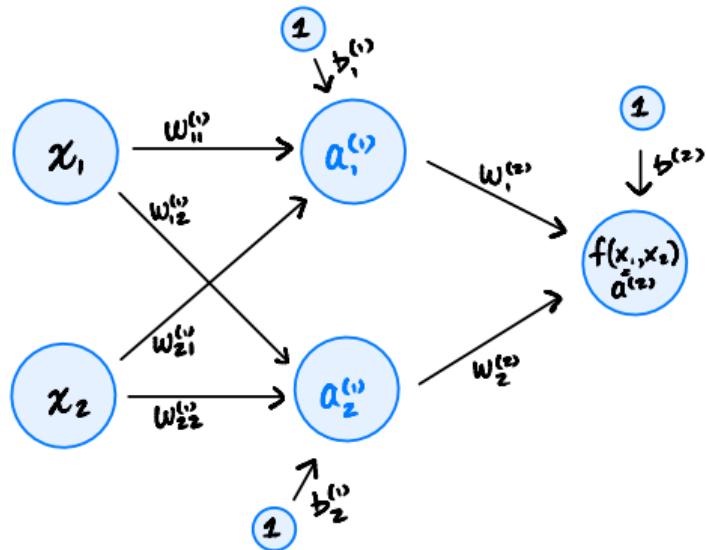
# DSC 190

## *Machine Learning: Representations*

Lecture 14 | Part 4

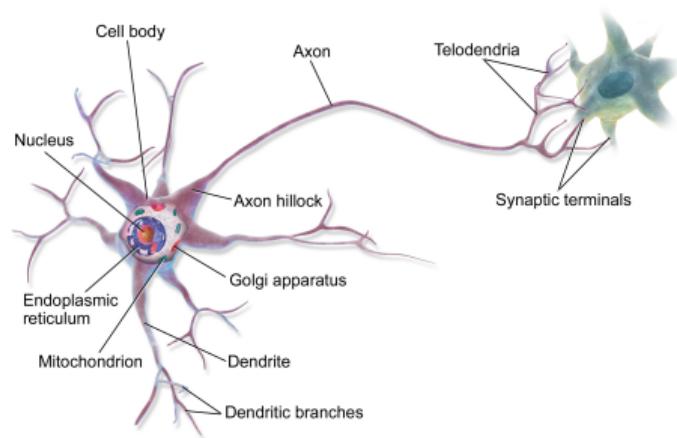
### **Hidden Units**

# Hidden Units



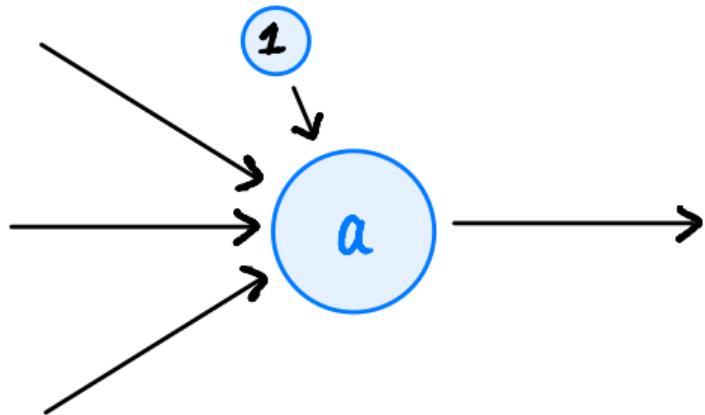
# Neuron

- ▶ Neuron accepts signals along **synapses**.
- ▶ Synapses have weights.
- ▶ If weighted sum is “large enough”, the neuron fires, or **activates**.



# Neuron

- ▶ Neuron accepts weighted inputs.
- ▶ If weighted sum is “large enough”, the neuron fires, or **activates**.

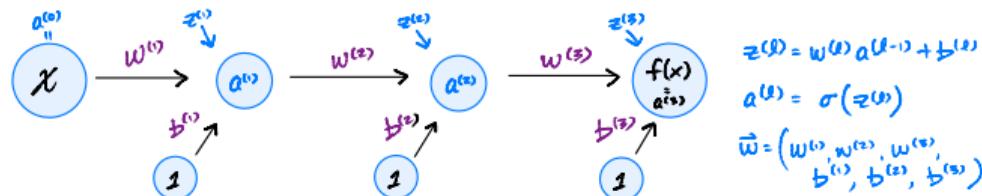


# Activation Functions

- ▶ A function  $g$  determining whether – and how strong – a neuron fires.
- ▶ We have seen two: ReLU and linear.
- ▶ Many different choices.
- ▶ Guided by intuition and only a little theory.

# Backpropagation

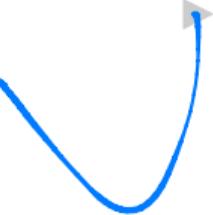
- ▶ The choice of activation function affects performance of backpropagation.
- ▶ Example:



$$\begin{aligned} z^{(l)} &= w^{(l)} a^{(l-1)} + b^{(l)} \\ a^{(l)} &= \sigma(z^{(l)}) \\ \vec{w} &= (w^{(1)}, w^{(2)}, w^{(3)}, b^{(1)}, b^{(2)}, b^{(3)}) \end{aligned}$$

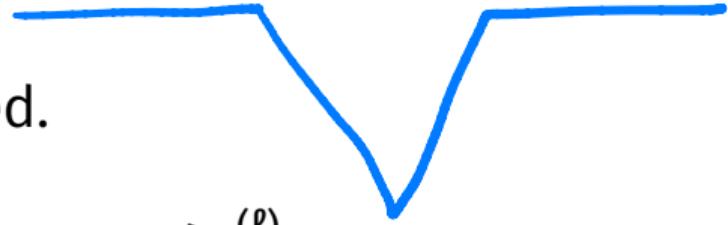
$$\frac{\partial f}{\partial w^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot \underbrace{g'(z^{(\ell)})}_{\sigma'} \cdot \frac{\partial z^{(\ell)}}{\partial w^{(\ell)}}$$

# Vanishing Gradients



A major challenge in training deep neural networks with backpropagation is that of **vanishing gradients**.

- ▶ The gradient for layers far from the output becomes very small.
- ▶ Weights can't be changed.


$$\frac{\partial f}{\partial w^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot g'(z^{(\ell)}) \cdot \frac{\partial z^{(\ell)}}{\partial w^{(\ell)}}$$

## Main Idea

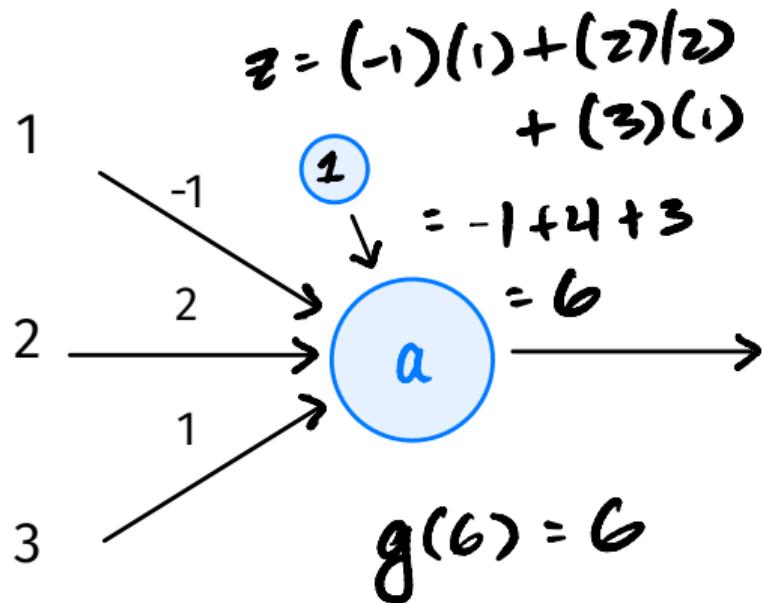
Some activation functions promote “healthier” gradients.

# Linear Activations

- A **linear** unit's activation function is:

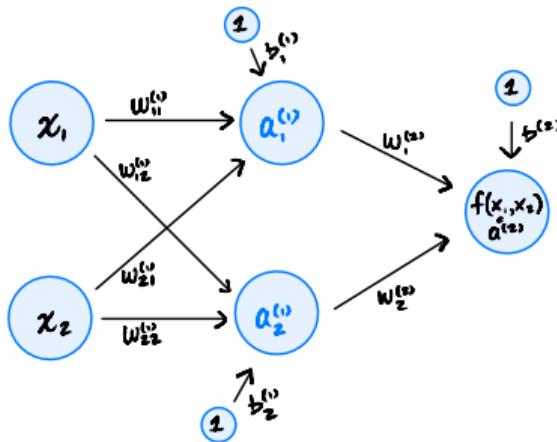
$$g(z) = z$$

$$\frac{\partial g}{\partial z} = 1$$

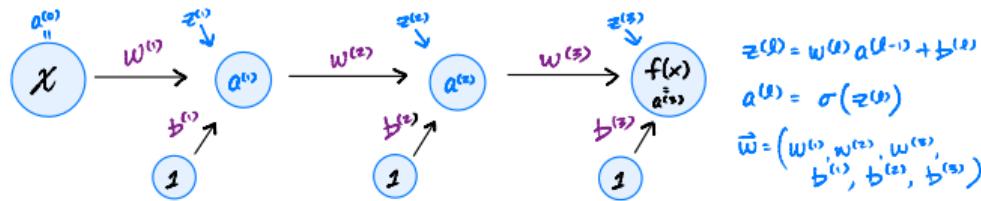


# Problem

- ▶ Linear activations result in a linear prediction function.



# Backprop. with Linear Activations



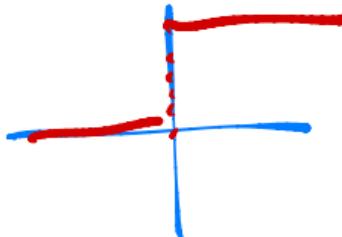
$$\frac{\partial f}{\partial w^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot g'(z^{(\ell)}) \cdot \frac{\partial z^{(\ell)}}{\partial w^{(\ell)}}$$

# Summary: Linear Activations

- ▶ **Good:** healthy gradients, fast to compute
- ▶ **Bad:** still results in linear prediction function when layers are combined

# Sigmoidal Activations

- ▶ A basic nonlinearity.
- ▶ Neuron is either “on” (1), “off” (0), or somewhere in between.
- ▶ Very popular before introduction of the ReLU.

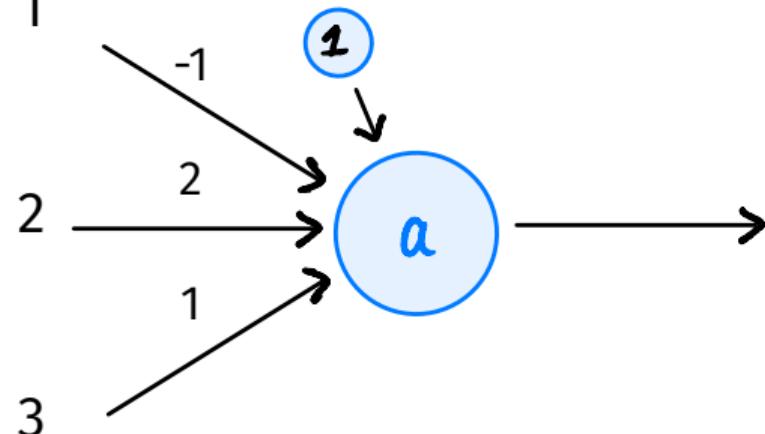
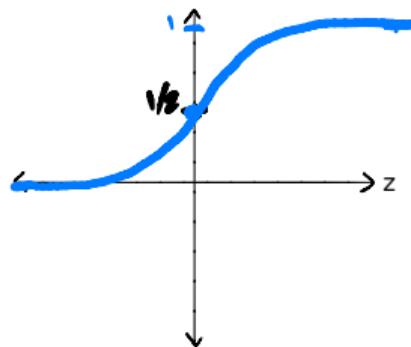


# Sigmoidal Activations

- A **sigmoidal** unit's activation function is:

$$g(z) = \frac{1}{1 + e^{-z}}$$

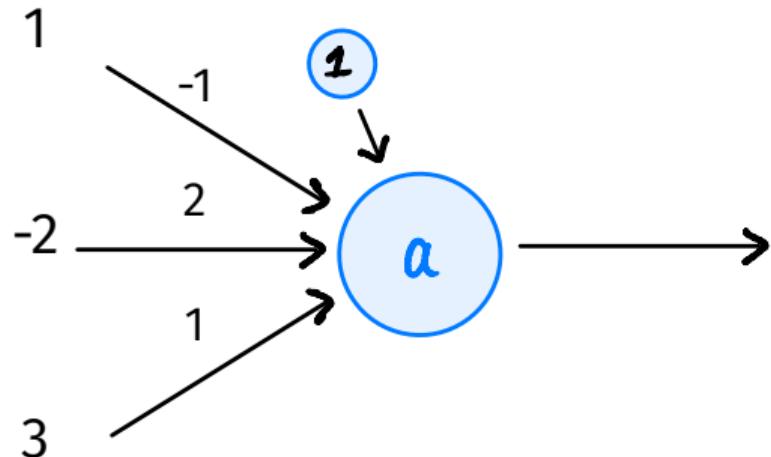
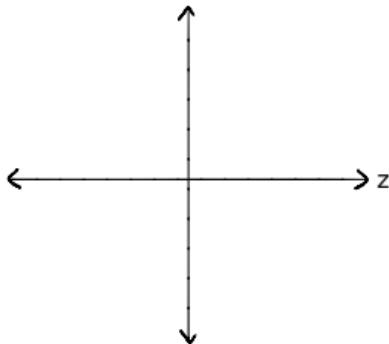
$$g(\infty) = \frac{1}{1+0} = 1 \quad g(-\infty) \approx 0$$



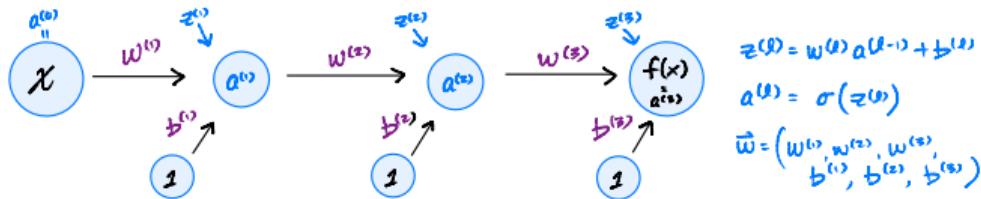
# Sigmoidal Activations

- A **sigmoidal** unit's activation function is:

$$g(z) = \frac{1}{1 + e^{-z}}$$



# Backprop. with Sigmoids



$$g'(z) = g(z)(1 - g(z)) \quad \frac{\partial f}{\partial w^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot g'(z^{(\ell)}) \cdot \frac{\partial z^{(\ell)}}{\partial w^{(\ell)}}$$

# Problem: Saturation

- ▶ Large/small inputs lead  $g(z)$  to be very close to 1 or -1.
- ▶ Here, the derivative  $\sigma'(z) \approx 0$ .
- ▶ Vanishing gradients!
- ▶ Makes learning deep networks with gradient-based algorithms very difficult.

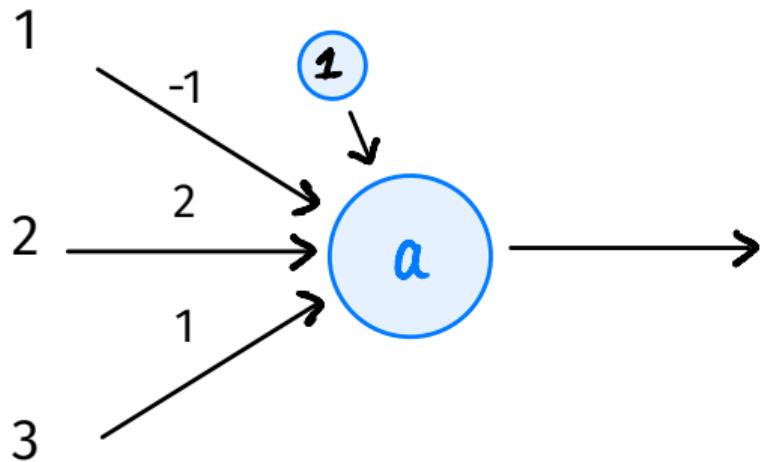
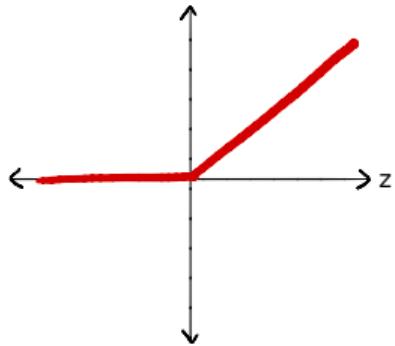
# ReLU

- ▶ Linear activations have strong gradients, but combined are still linear.
- ▶ Sigmoidal activations are non-linear, but when saturated lead to weak gradients.
- ▶ Can we have the best of both?

# ReLU

- A **rectified linear** unit's (ReLU) activation function is:

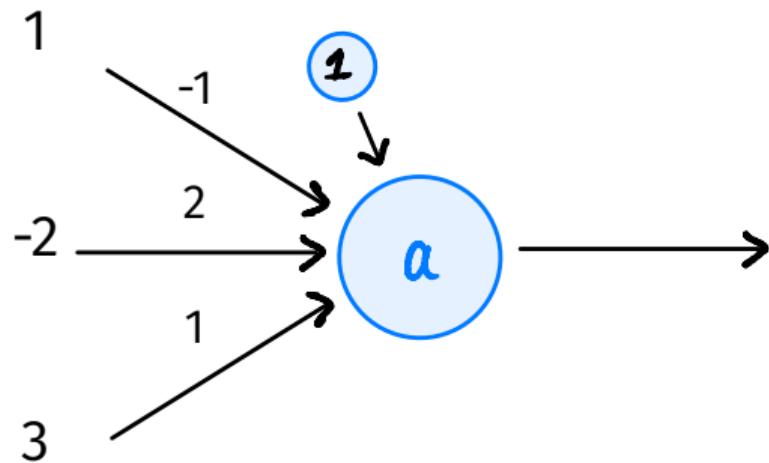
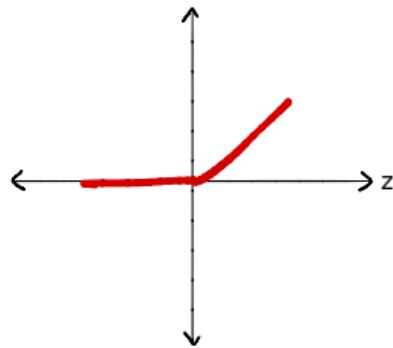
$$g(z) = \max\{0, z\}$$



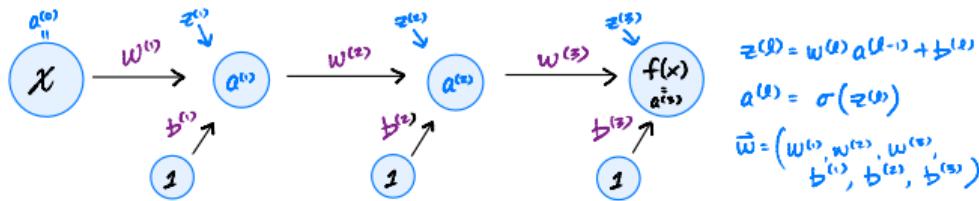
# ReLU

- A **rectified linear** unit's (ReLU) activation function is:

$$g(z) = \max\{0, z\}$$



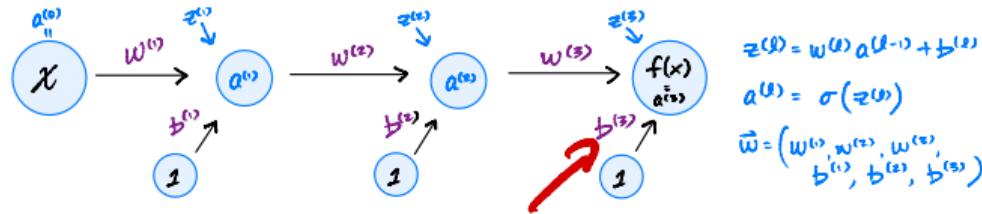
# Backprop. with ReLU



$$\frac{\partial f}{\partial w^{(\ell)}} = \frac{\partial f}{\partial a^{(\ell)}} \cdot g'(z^{(\ell)}) \cdot \frac{\partial z^{(\ell)}}{\partial w^{(\ell)}}$$

# Backprop. with ReLU

- ▶ **Problem:** If inputs < 0, ReLU “deactivates” and gradients are not passed back.



-10

# Fixing Deactivated ReLUs

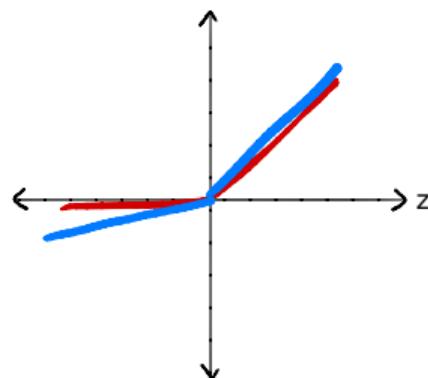
- ▶ One fix: initialize all biases to be small, positive numbers.
- ▶ Ensures that most units are active to begin with.
- ▶ Another fix: modify the ReLU.

# Leaky ReLU

- ▶ A **leaky ReLU** activation function is:

$$g(z) = \max\{\alpha z, z\} \quad 0 \leq \alpha < 1$$

- ▶ Usually,  $\alpha \approx 0.01$ . Nonzero derivative.



# Summary: ReLU

- ▶ The popular, “default” choice of activation function.
- ▶ **Good:** Strong gradient when active, fast to compute.
- ▶ **Bad:** No gradient when inactive.

# DSC 190

## Machine Learning: Representations

Lecture 14 | Part 5

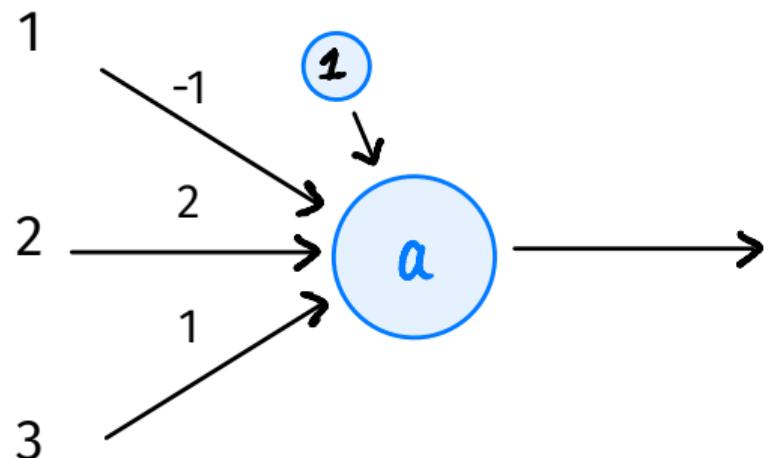
### Output Units

# Output Units

- ▶ As with units in hidden layers, we can customize output units.
  - ▶ What activation function?
  - ▶ How many units?
- ▶ Good choice depends on task:
  - ▶ Regression, binary classification, multiclass, etc.
- ▶ Which loss?

# Setting 1: Regression

- ▶ Output can be any real number.
- ▶ Single output neuron.
- ▶ It makes sense to use a **linear activation**.

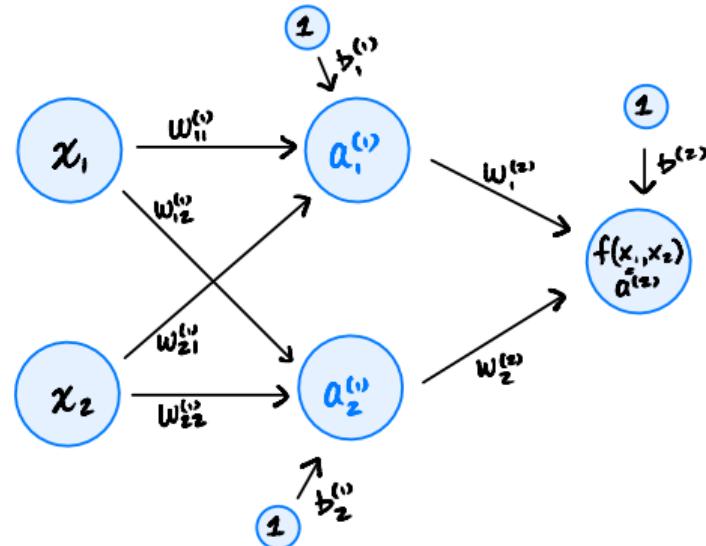


# Setting 1: Regression

- ▶ Prediction should not be too high/low.
- ▶ It makes sense to use the **mean squared error**.

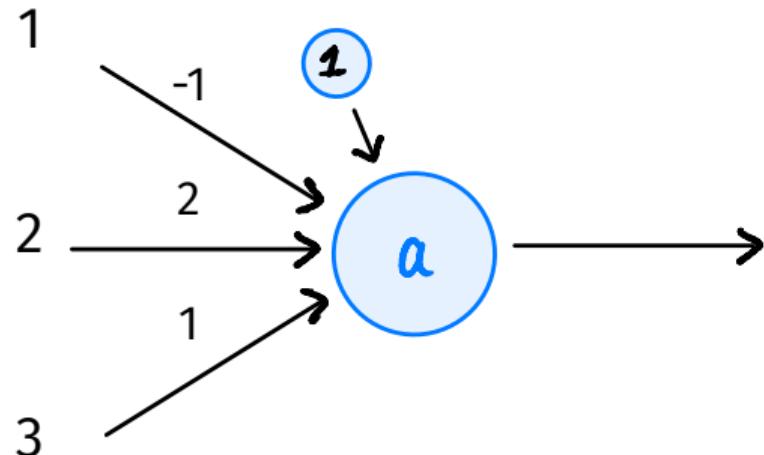
# Setting 1: Regression

- ▶ Suppose we use linear activation for output neuron + mean squared error.
- ▶ This is very similar to least squares regression...
- ▶ But! Features in earlier layers are **learned**, non-linear.



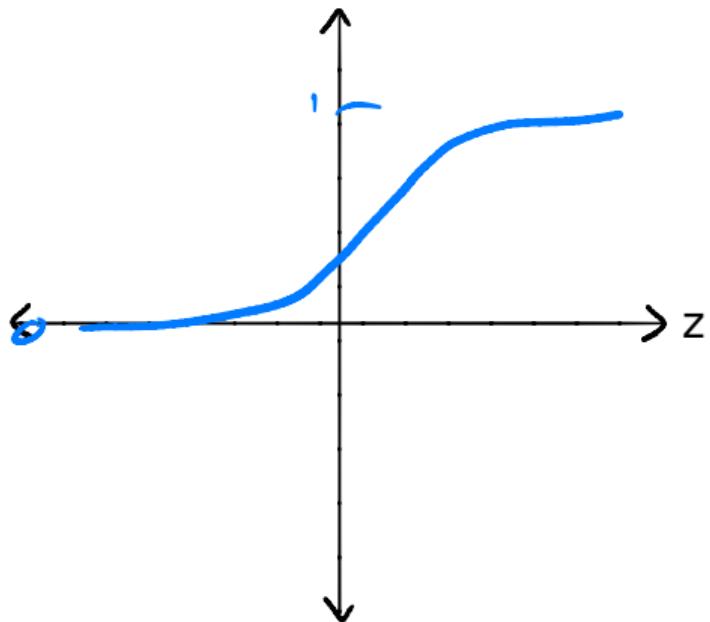
# Setting 2: Binary Classification

- ▶ Output can be in  $[0, 1]$ .
- ▶ Single output neuron.
- ▶ We *could* use a **linear activation**, threshold.
- ▶ But there is a better way.



# Sigmoids for Classification

- ▶ Natural choice for activation in output layer for binary classification: the **sigmoid**.



# Binary Classification Loss

- ▶ We could use square loss for binary classification. There are several reasons not to:
- ▶ 1) Square loss penalizes predictions which are “too correct”.
- ▶ 2) It doesn’t work well with the sigmoid due to saturation.

# The Cross-Entropy

- ▶ Instead, we often train deep classifiers using the **cross-entropy** as loss.
- ▶ Let  $y^{(i)} \in \{0, 1\}$  be true label of  $i$ th example.
- ▶ The average cross-entropy loss:

$$-\frac{1}{n} \sum_{i=1}^n \begin{cases} \log f(\vec{x}^{(i)}), & \text{if } y^{(i)} = 1 \\ \log[1 - f(\vec{x}^{(i)})], & \text{if } y^{(i)} = 0 \end{cases}$$

# The Cross-Entropy and the Sigmoid

- ▶ Cross-entropy “undoes” the exponential in the sigmoid, resulting in less saturation.

# Summary: Binary Classification

- ▶ Use sigmoidal activation the output layer + cross-entropy loss.
- ▶ This will promote a strong gradient.
- ▶ Use whatever activation for the hidden layers (e.g., ReLU).