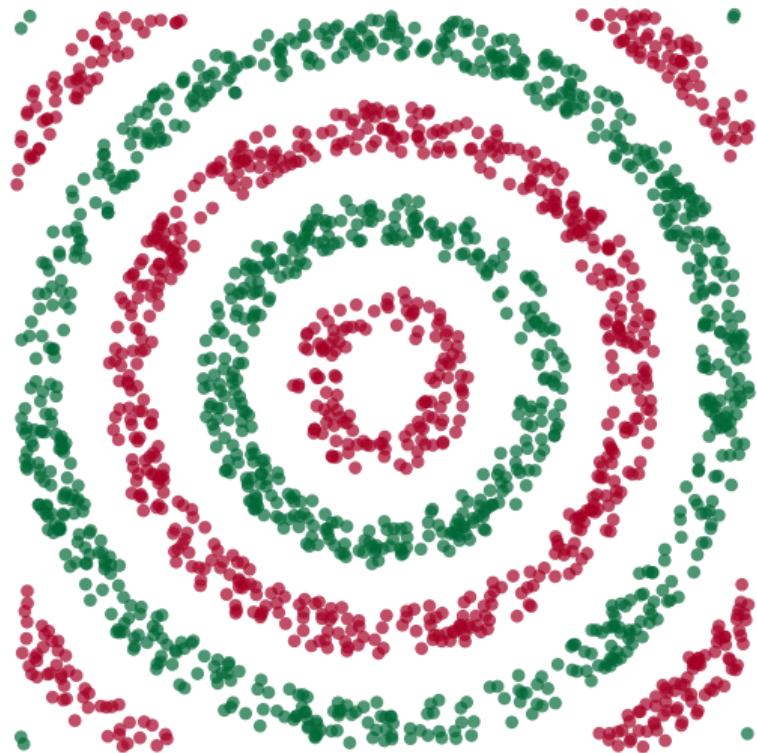


DSC 190

Machine Learning: Representations

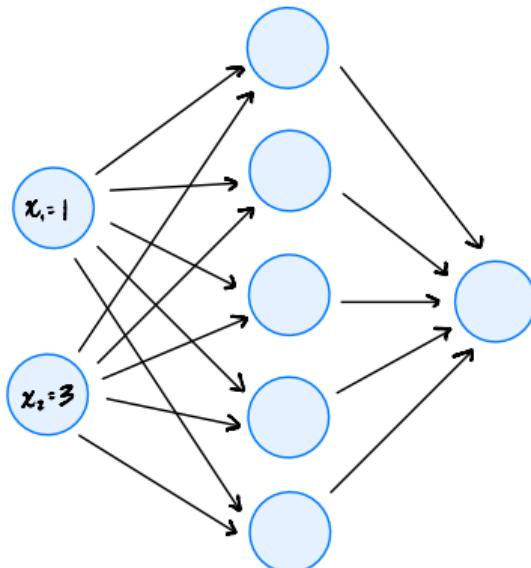
Lecture 15 | Part 1

NNs and Representations



NNs and Representations

- ▶ Hidden layer transforms to new **representation**.
 - ▶ Maps $\mathbb{R}^2 \rightarrow \mathbb{R}^5$
- ▶ Output layer makes prediction.
 - ▶ Maps $\mathbb{R}^5 \rightarrow \mathbb{R}^1$
- ▶ Representation optimized for classification!



NN Design

- ▶ Design a network for classification.
- ▶ Hidden layer activations: ReLU
- ▶ Output layer activation: sigmoid
- ▶ Loss function: cross-entropy

```
from tensorflow import keras

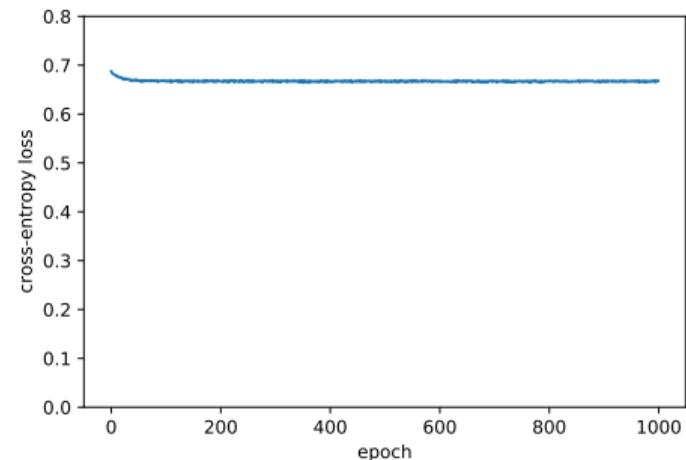
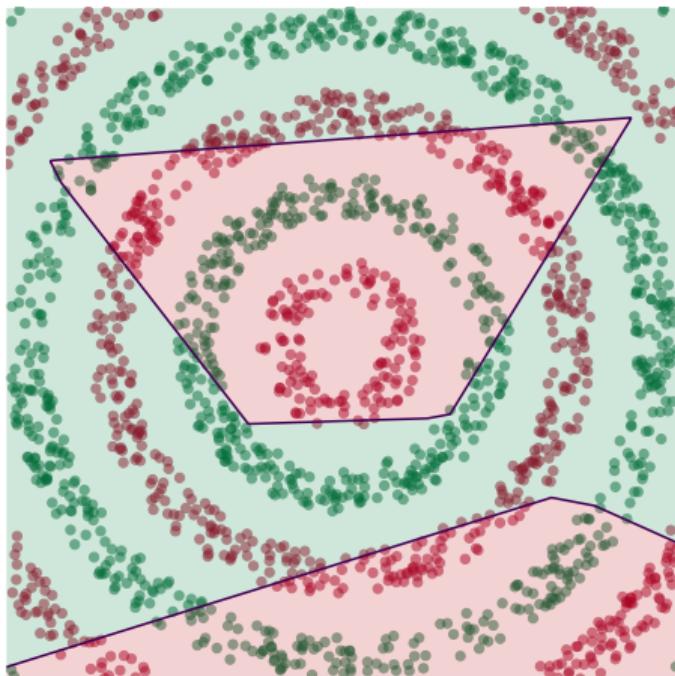
inputs = keras.Input(shape=2)
hidden_1 = keras.layers.Dense(5, activation='relu')(inputs)
outputs = keras.layers.Dense(1, activation='sigmoid')(hidden_1)

model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(
    optimizer=keras.optimizers.RMSprop(learning_rate=.01),
    loss=keras.losses.BinaryCrossentropy()
)

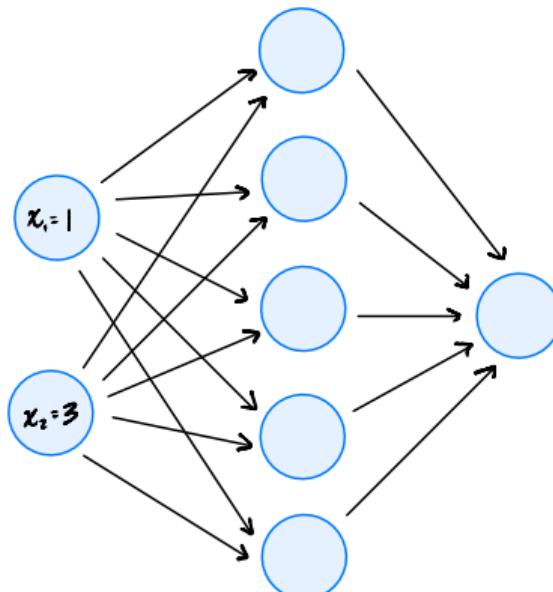
history = model.fit(X, y, epochs=1000, verbose=1)
```

Results



NNs and Representations

- ▶ Data has complex structure.
- ▶ Only 5 hidden neurons not enough to learn a good representation.



DSC 190

Machine Learning: Representations

Lecture 15 | Part 2

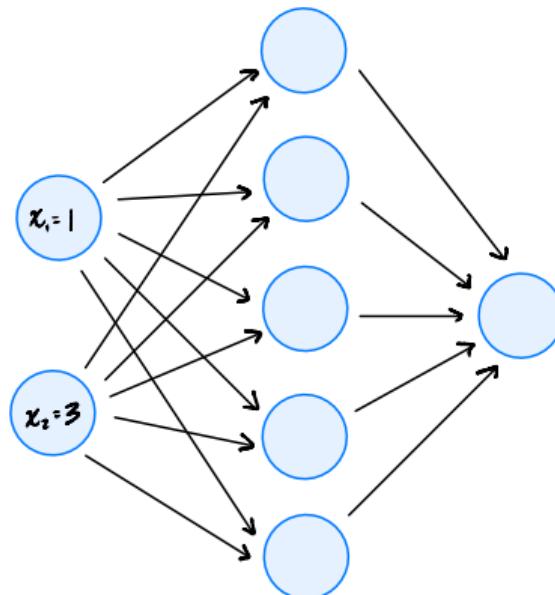
Architecture

Architecture

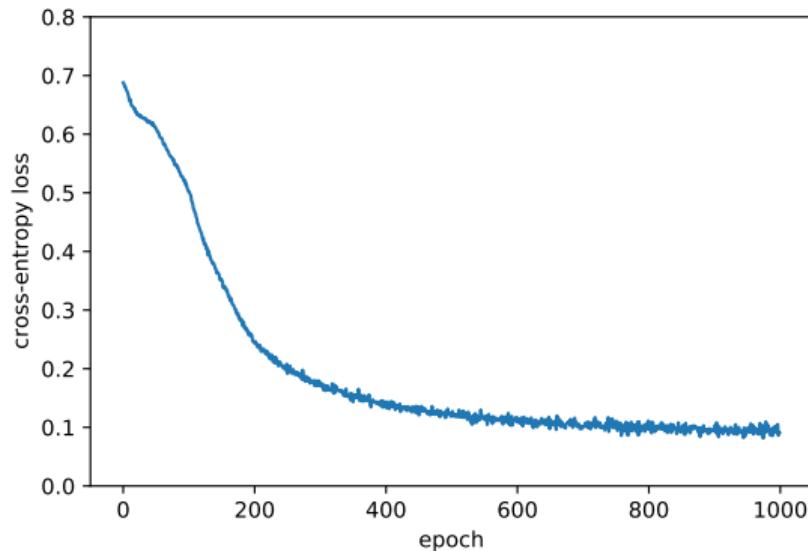
- ▶ We can increase complexity in two ways:
- ▶ Increasing **width**.
- ▶ Increasing **depth**.

Increasing Width

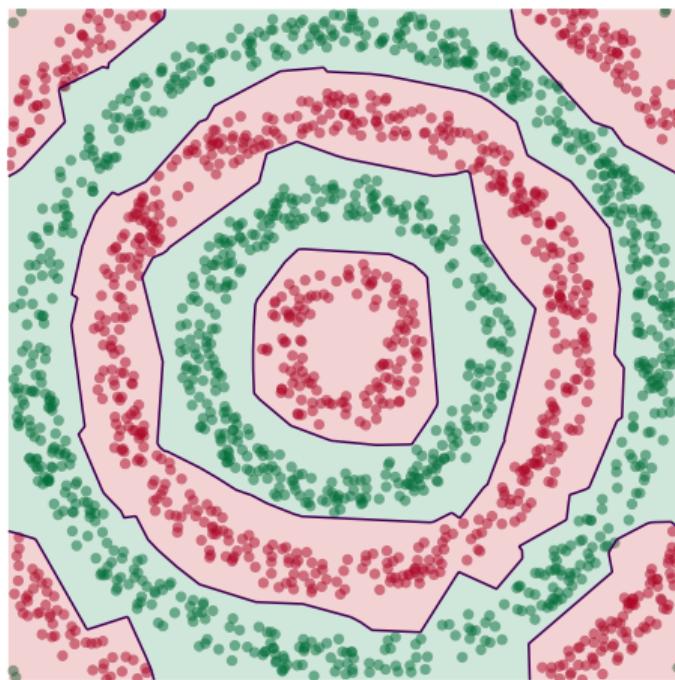
- ▶ Use a single hidden layer.
- ▶ But with 50 hidden neurons instead of 5.
- ▶ I.e., map to \mathbb{R}^{50} , then predict.



LOSS



Result



Universal Approximation Theorem

- ▶ A neural network f is a special type of function.
- ▶ Given another function g , can we make a neural network f so that $f(\vec{x}) \approx g(\vec{x})$?
- ▶ **Yes!** Assuming:
 - ▶ f has a hidden layer with a suitable activation function (ReLU, sigmoid, etc.)
 - ▶ the hidden layer has **enough** neurons
 - ▶ g is not too wild.

Main Idea

A network with a single hidden layer is able to approximate any (not-too-wild) function arbitrarily well as long as it has enough neurons in the hidden layer.

So what?

- ▶ Nature uses some function g to assign class labels to data.
- ▶ We don't see this function. But we see $g(\vec{x})$ for a bunch of points.
- ▶ Our goal is to learn a function f approximating g using this data.

The Challenge

- ▶ NNs are universal approximators (so are RBF networks, etc.)
- ▶ But just because it *can* approximate any function, doesn't mean we can *learn* the approximation.

Number of Neurons

- ▶ UAT says one hidden layer works well with “enough neurons”
- ▶ What is enough?
- ▶ Unfortunately, it can be a lot!

DSC 190

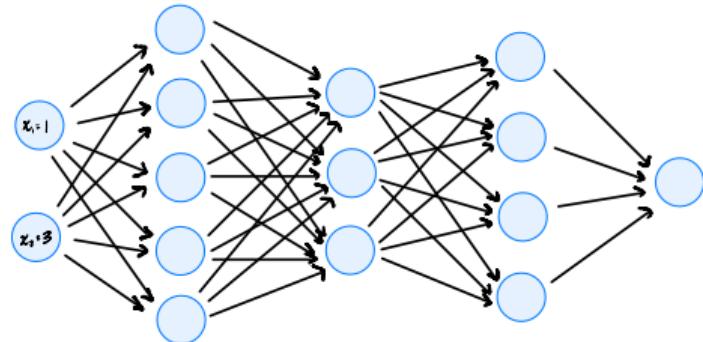
Machine Learning: Representations

Lecture 15 | Part 3

Deep Networks

Deep Networks

- ▶ Use a **multiple** hidden layers.
- ▶ Hidden layers transform to a new representation.
- ▶ Composition of simple transformations.
- ▶ Output layer performs prediction.



Main Idea

In machine learning, “deep” means “more than one hidden layer”. Deep models are useful for **learning** simpler representations.

Designing a Deep NN

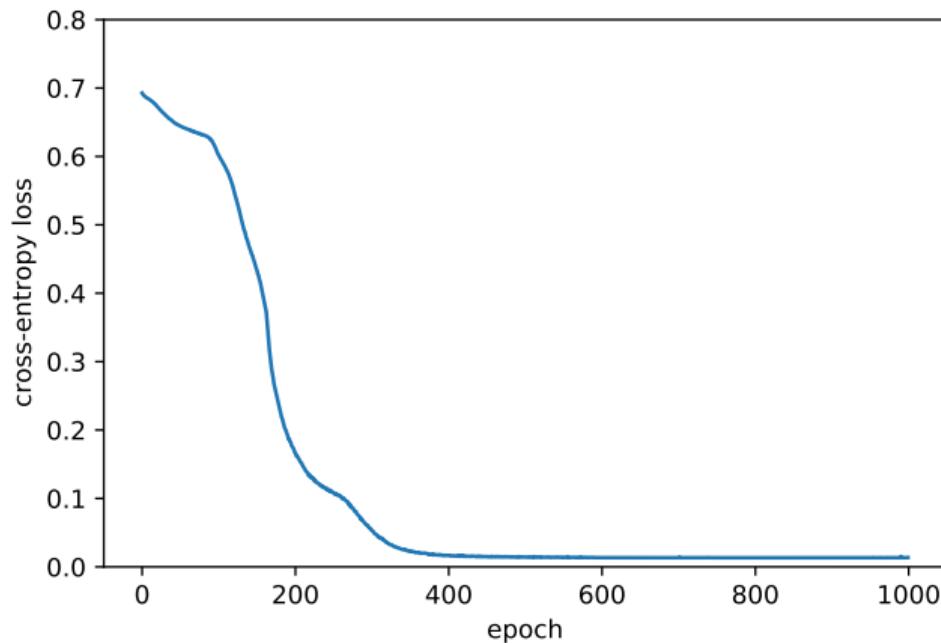
- ▶ Pick a number of hidden layers.
- ▶ Pick width of each hidden layer.
- ▶ There's not much theory to help us here.
- ▶ Experiment with different choices.

```
inputs = keras.Input(shape=2)
hidden_1 = keras.layers.Dense(15, activation='relu')(inputs)
hidden_2 = keras.layers.Dense(20, activation='relu')(hidden_1)
hidden_3 = keras.layers.Dense(2, activation='relu')(hidden_2)
outputs = keras.layers.Dense(1, activation='sigmoid')(hidden_3)

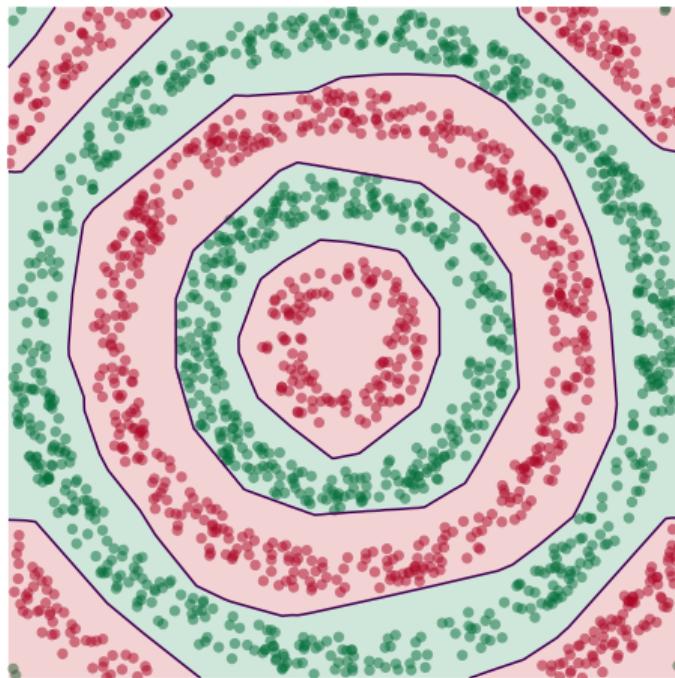
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(
    optimizer=keras.optimizers.RMSprop(learning_rate=.001),
    loss=keras.losses.BinaryCrossentropy()
)
history = model.fit(X, y, epochs=1000, verbose=1)
```

LOSS

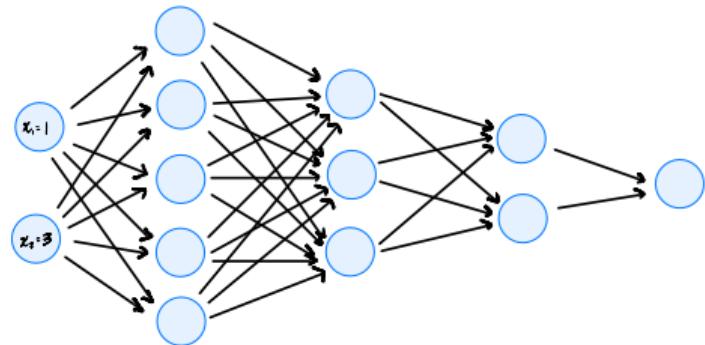


Result

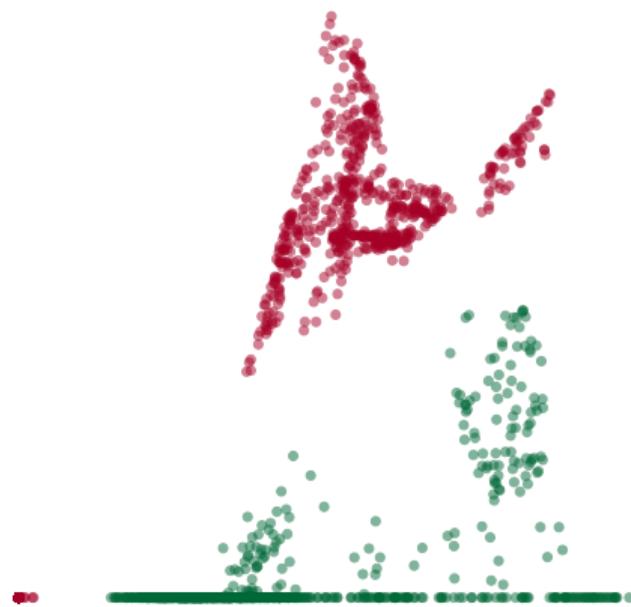


Deep Networks

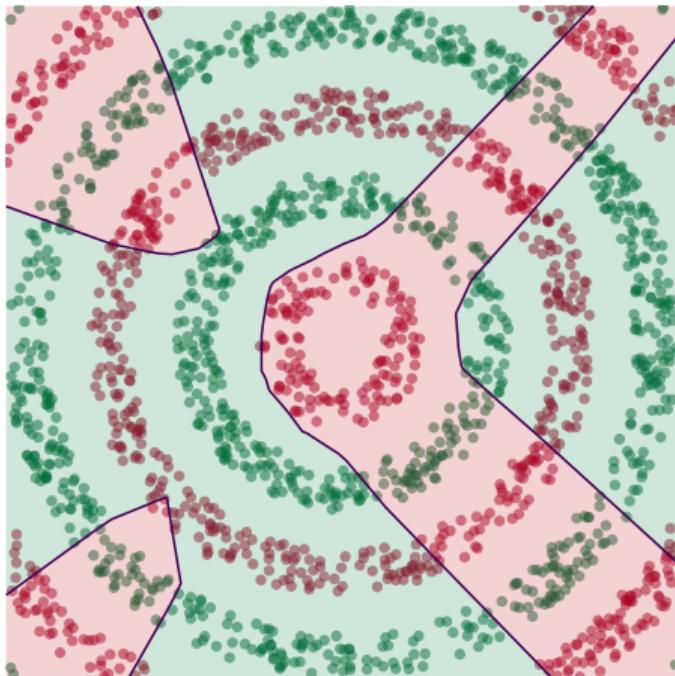
- ▶ Hidden layers map input to new representation.
- ▶ We can see this new representation!
- ▶ Plug in \vec{x} and see activations of last hidden layer.



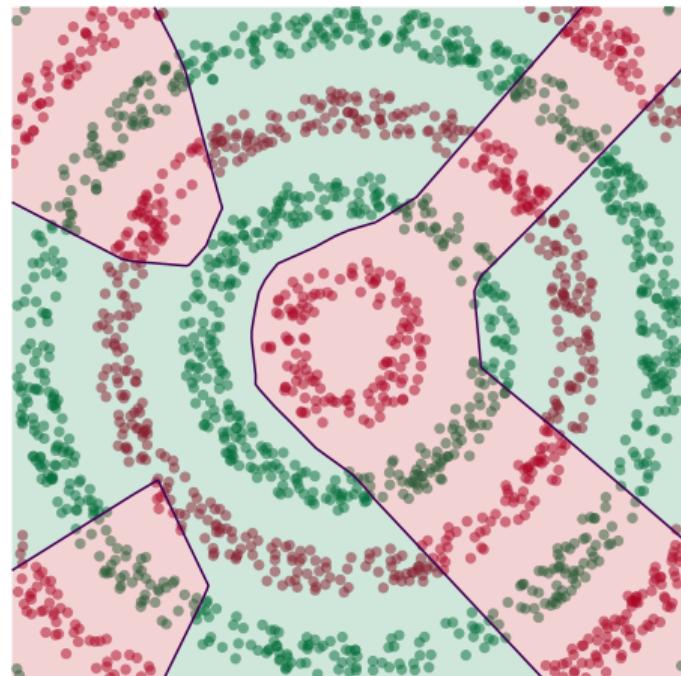
The New Representation



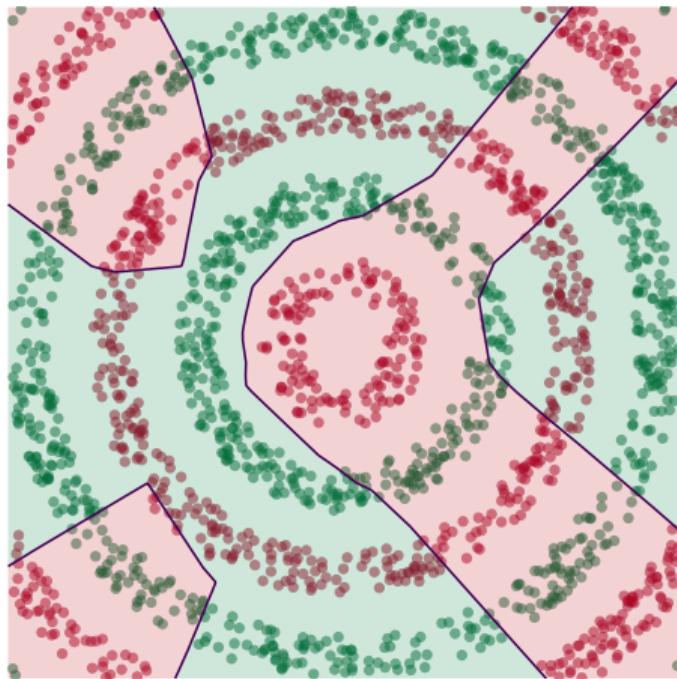
Learning a New Representation



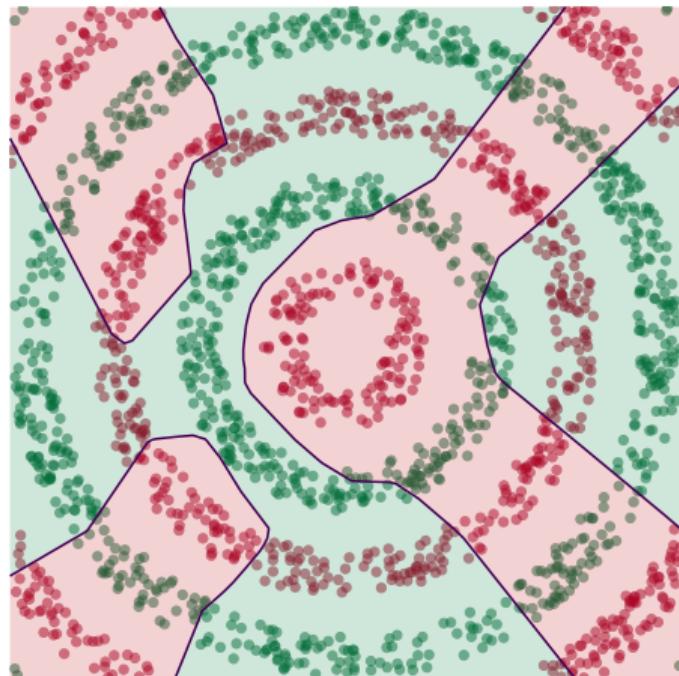
Learning a New Representation



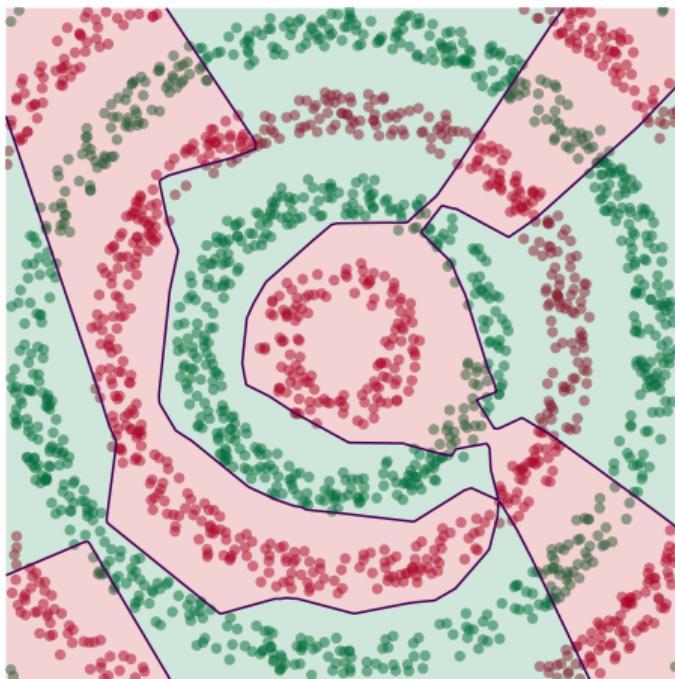
Learning a New Representation



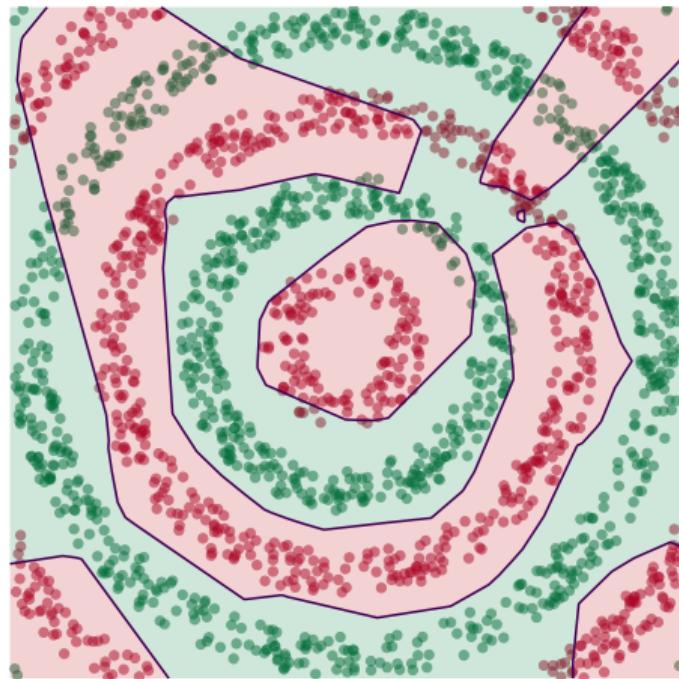
Learning a New Representation



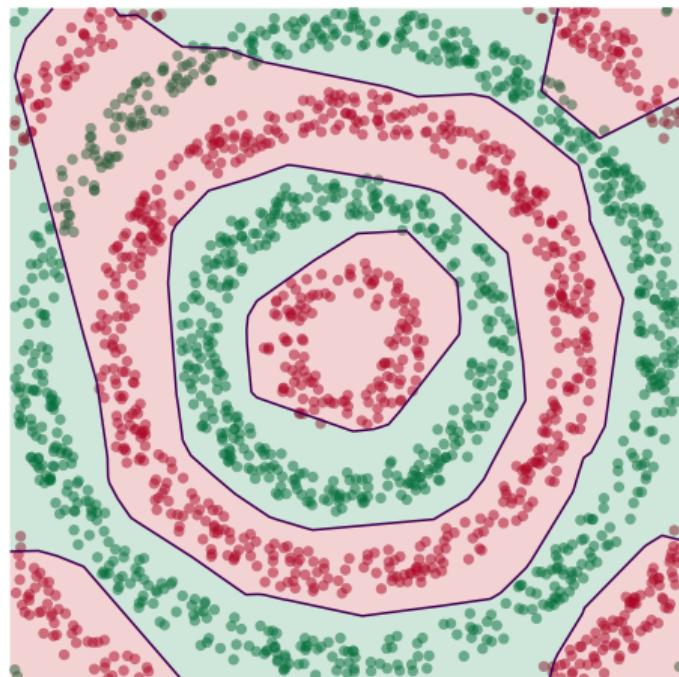
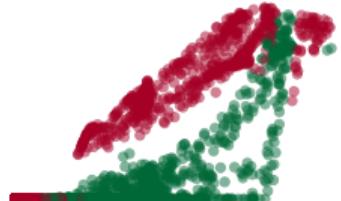
Learning a New Representation



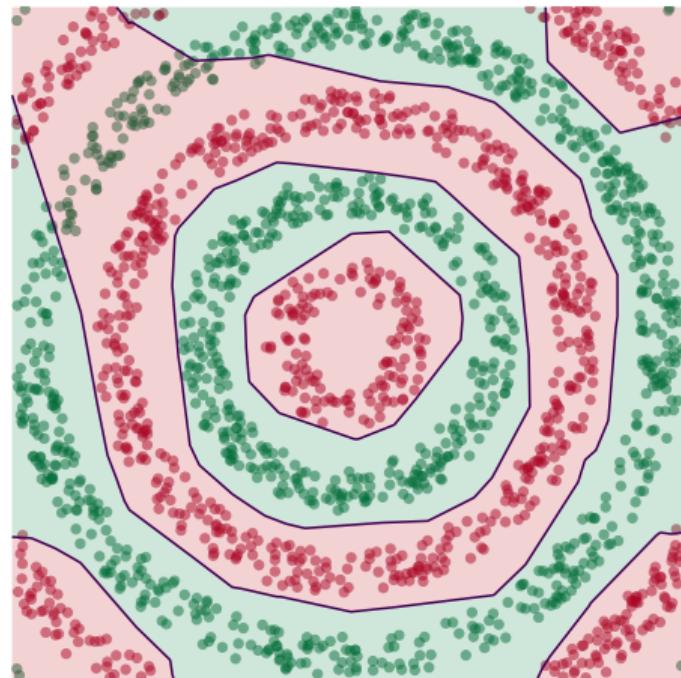
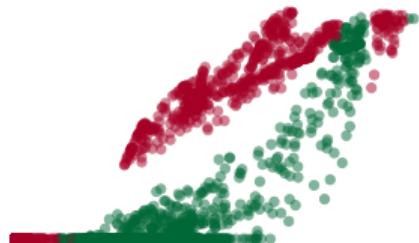
Learning a New Representation



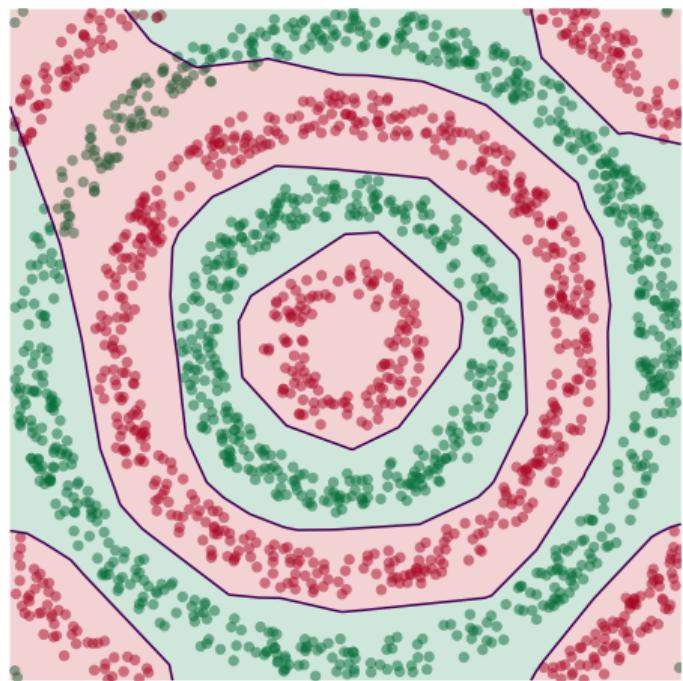
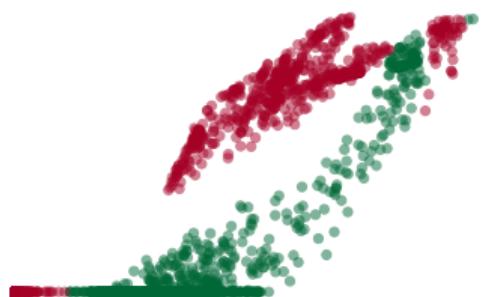
Learning a New Representation



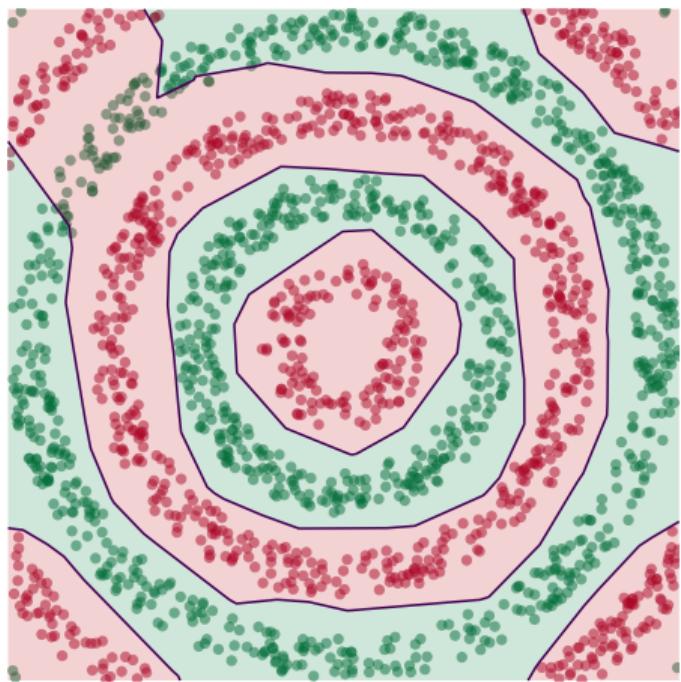
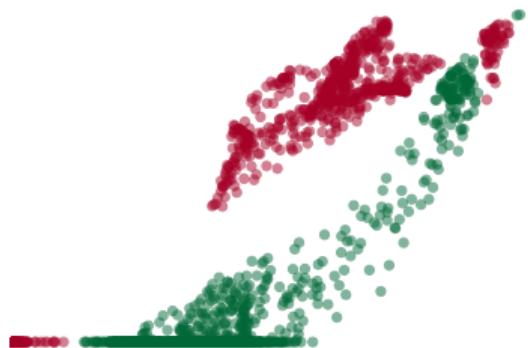
Learning a New Representation



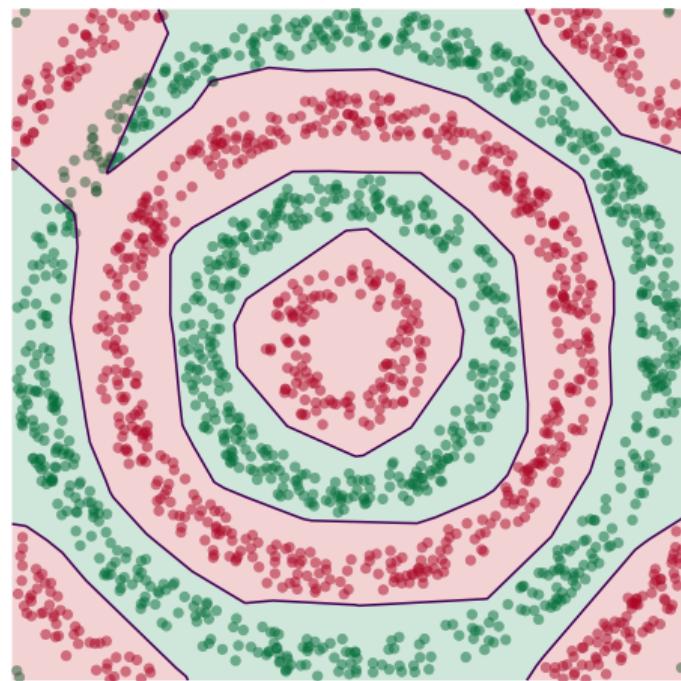
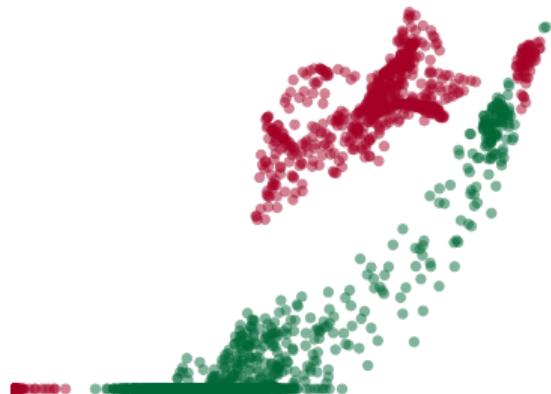
Learning a New Representation



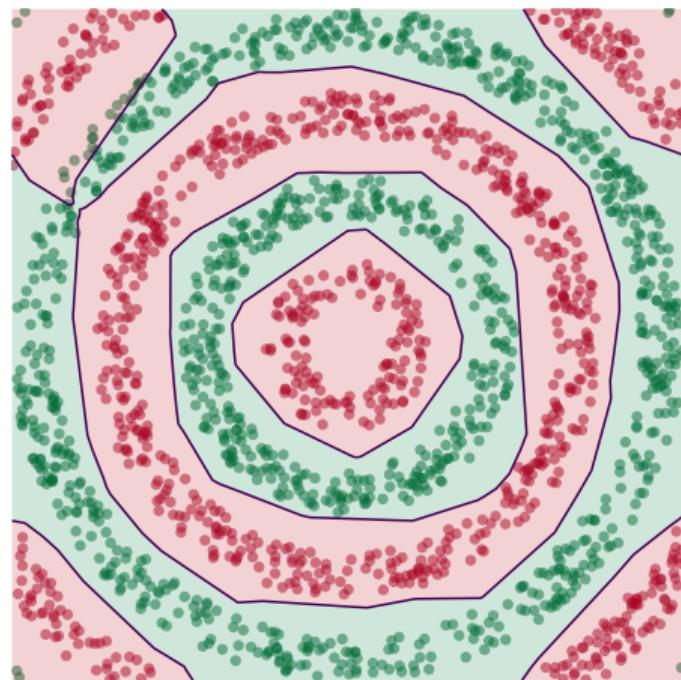
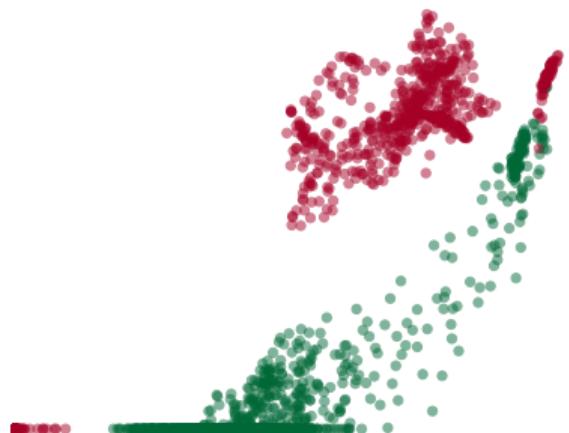
Learning a New Representation



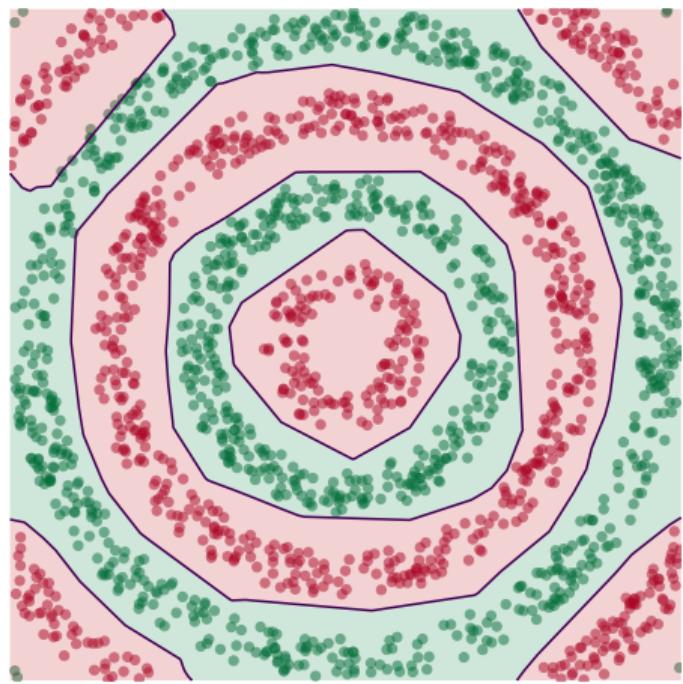
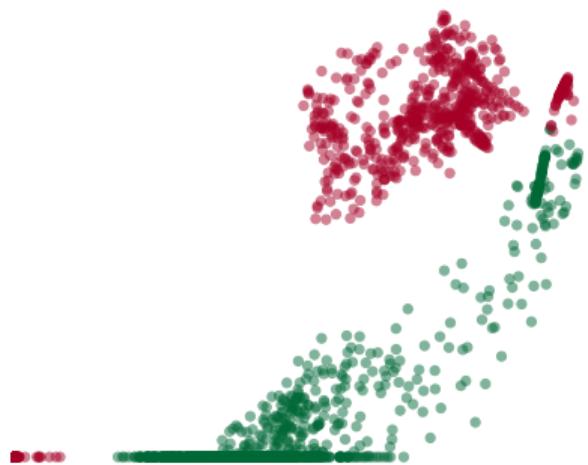
Learning a New Representation



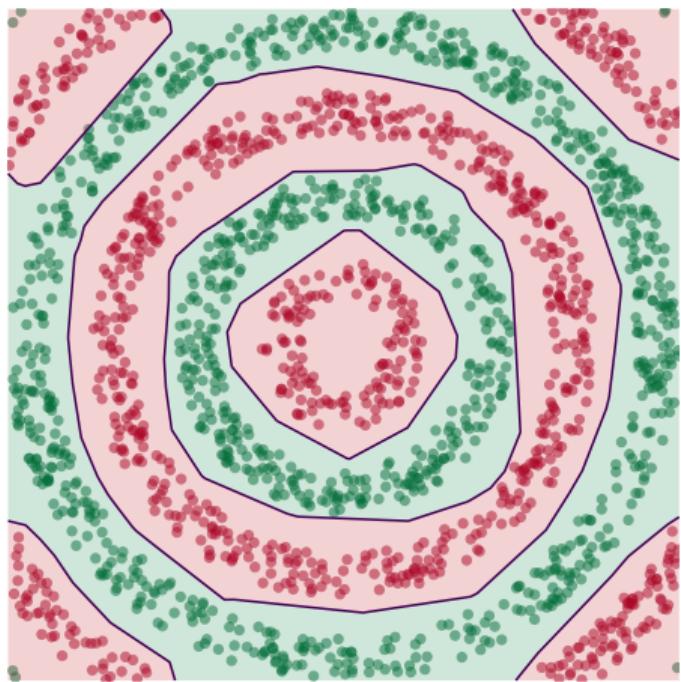
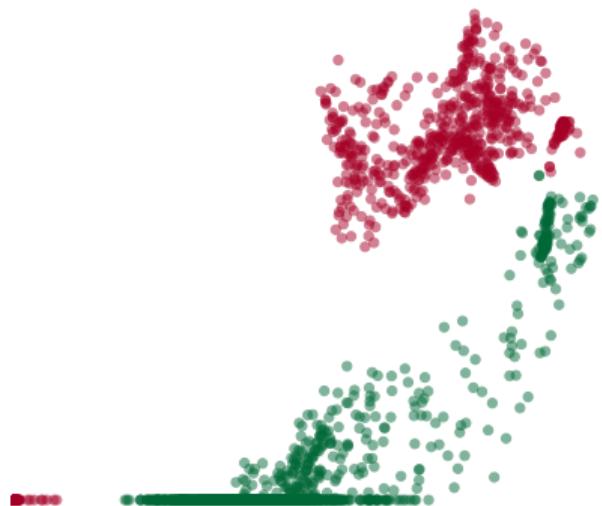
Learning a New Representation



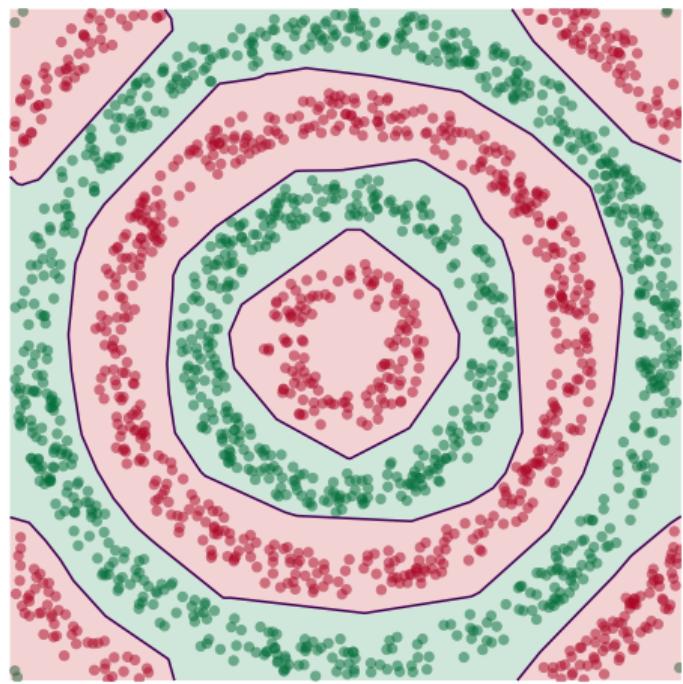
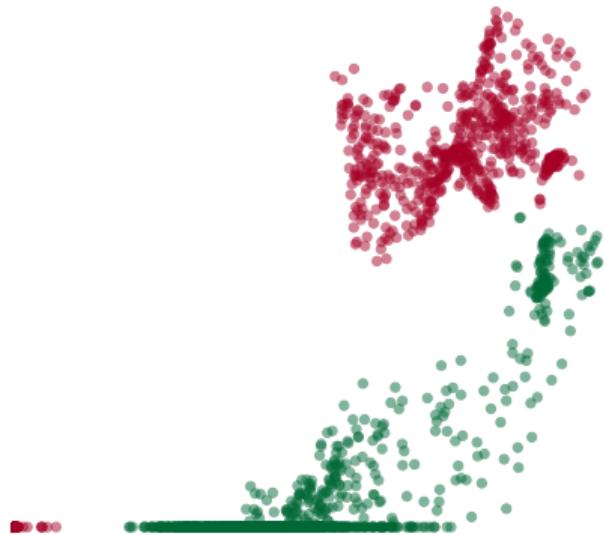
Learning a New Representation



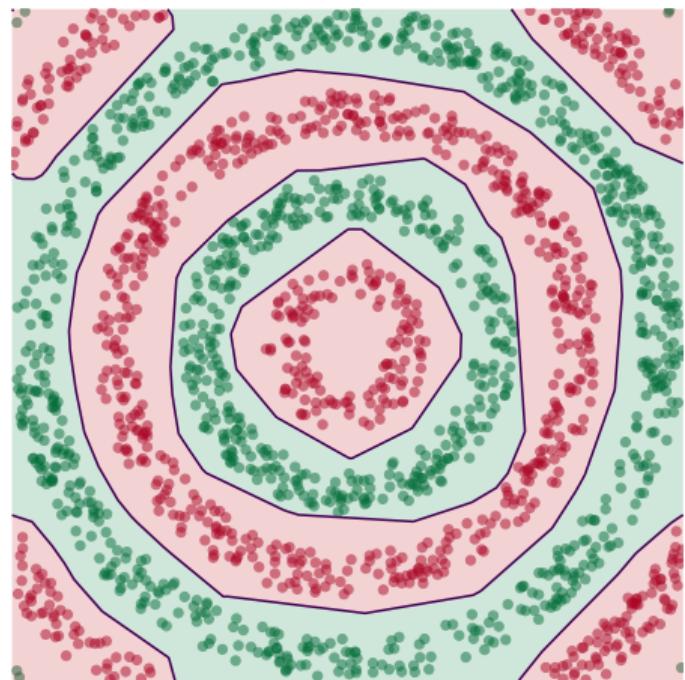
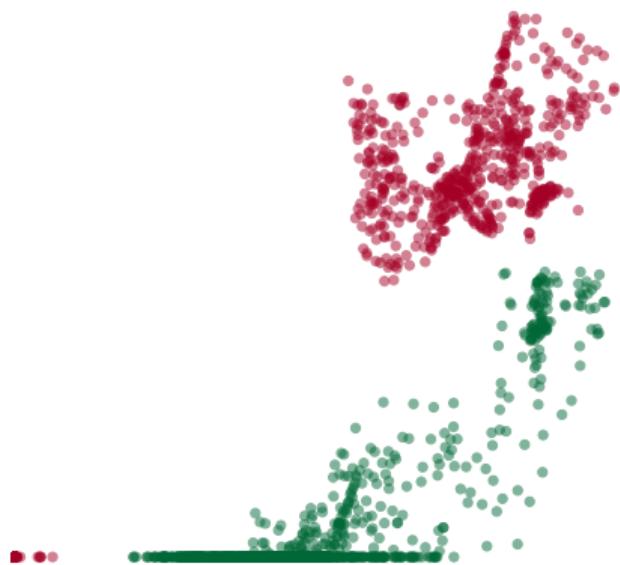
Learning a New Representation



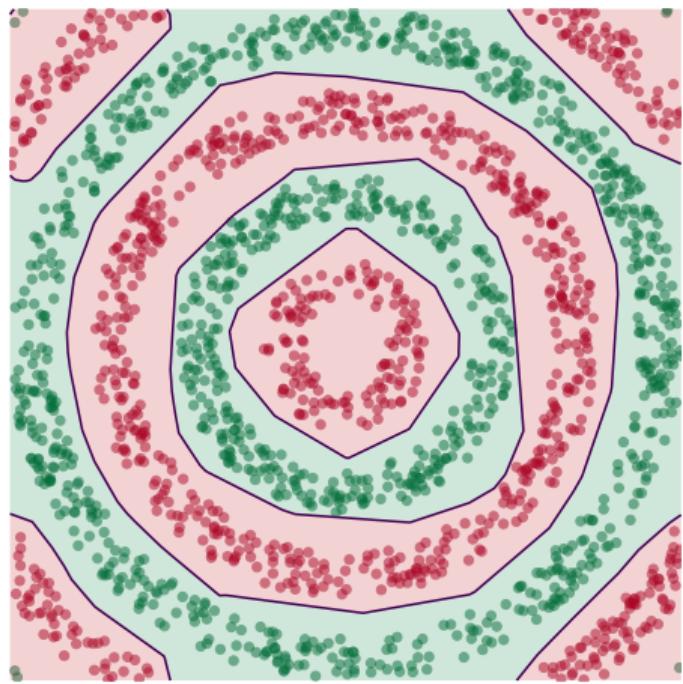
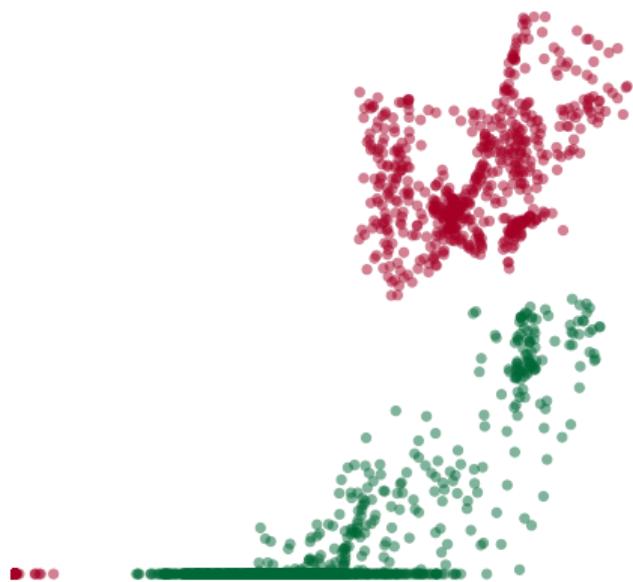
Learning a New Representation



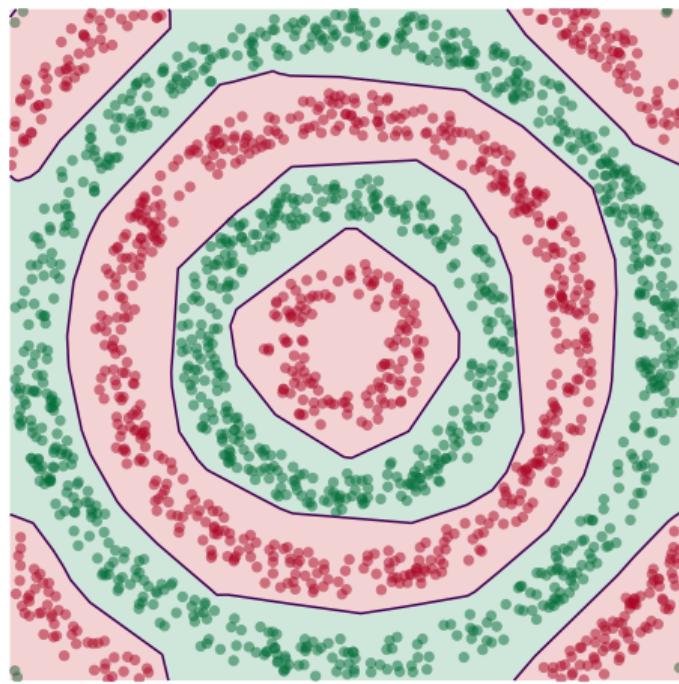
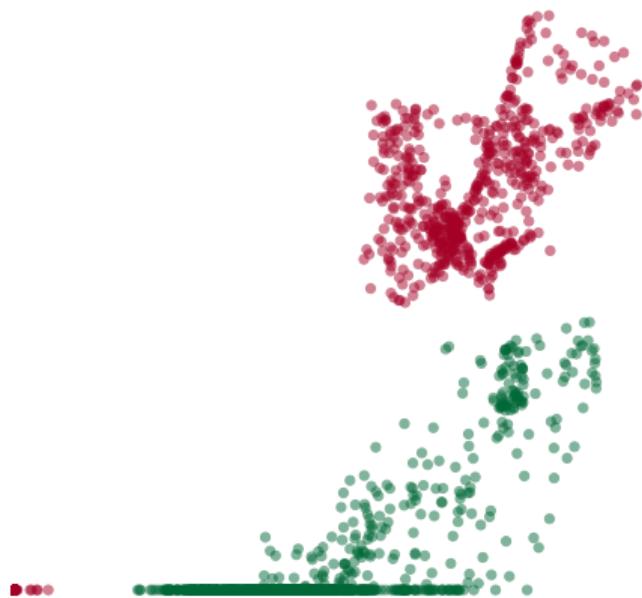
Learning a New Representation



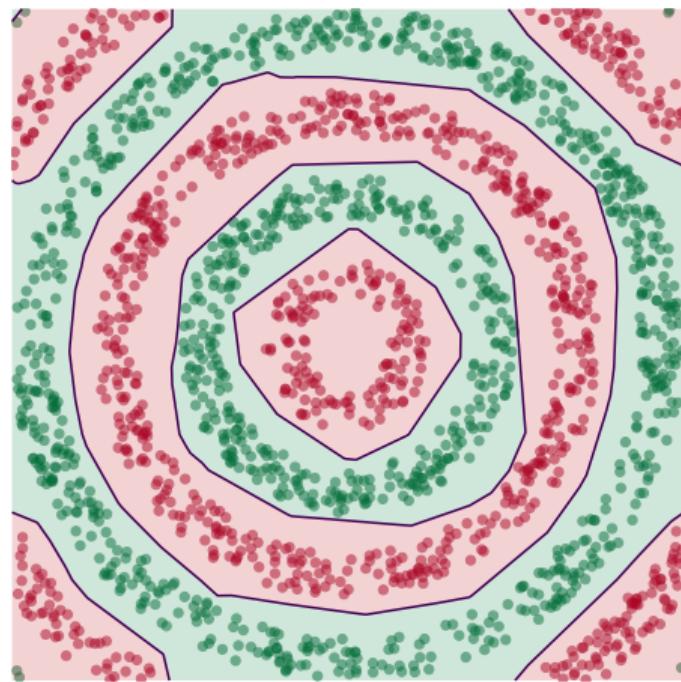
Learning a New Representation



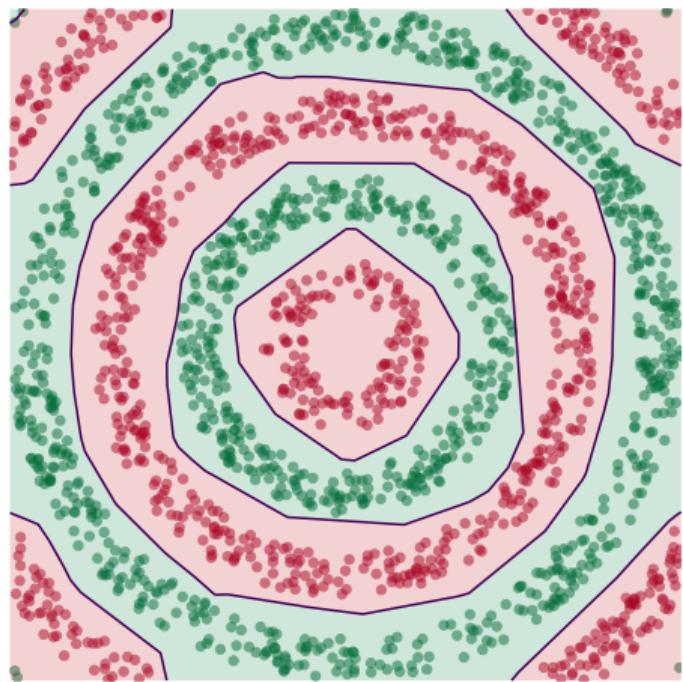
Learning a New Representation



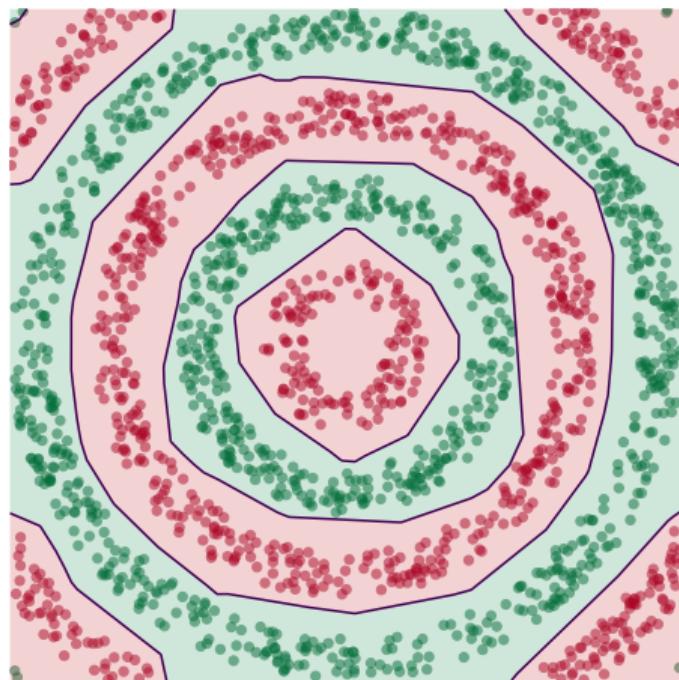
Learning a New Representation



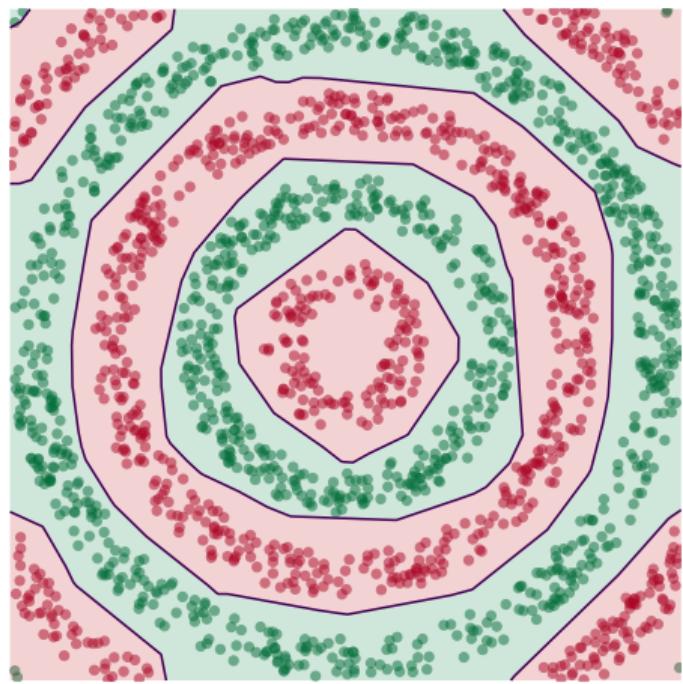
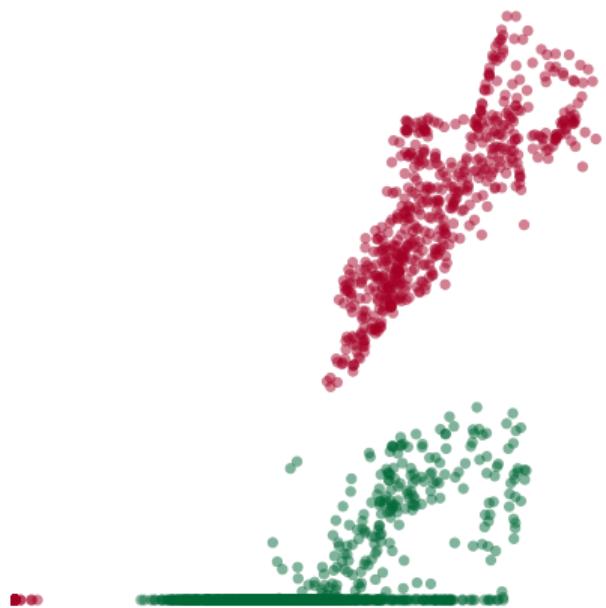
Learning a New Representation



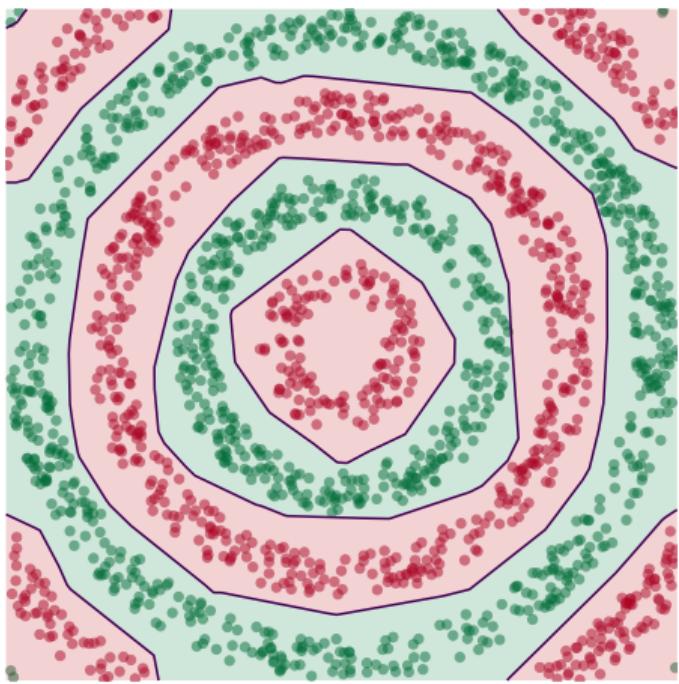
Learning a New Representation



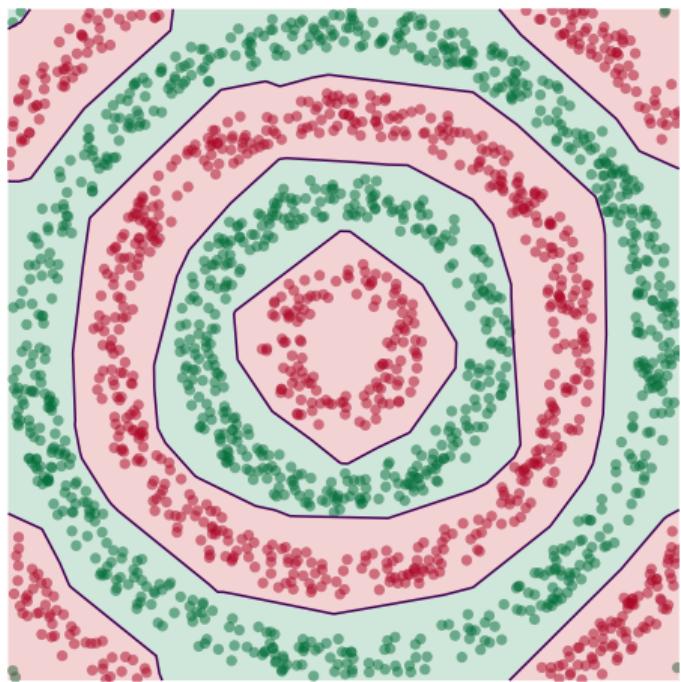
Learning a New Representation



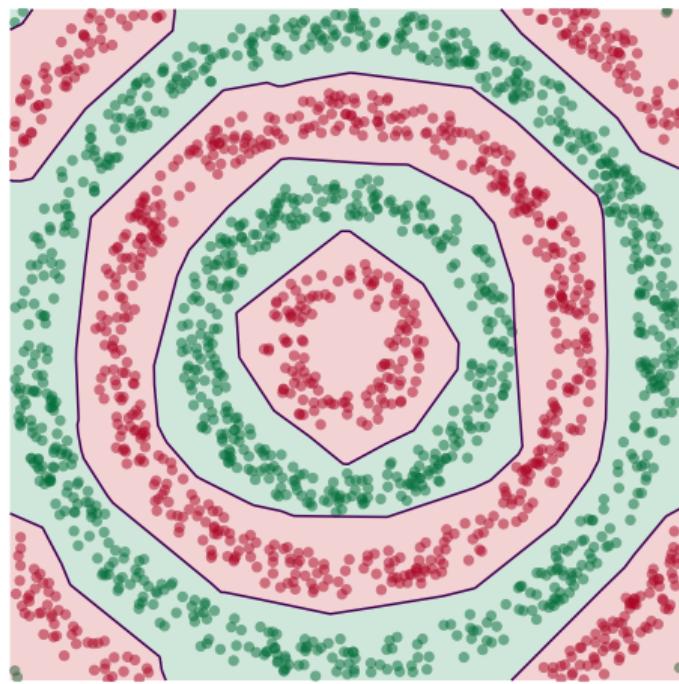
Learning a New Representation



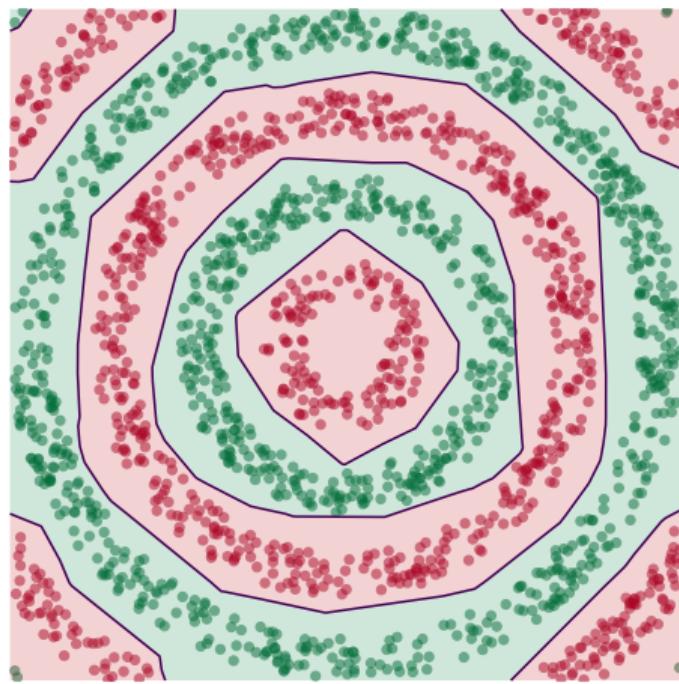
Learning a New Representation



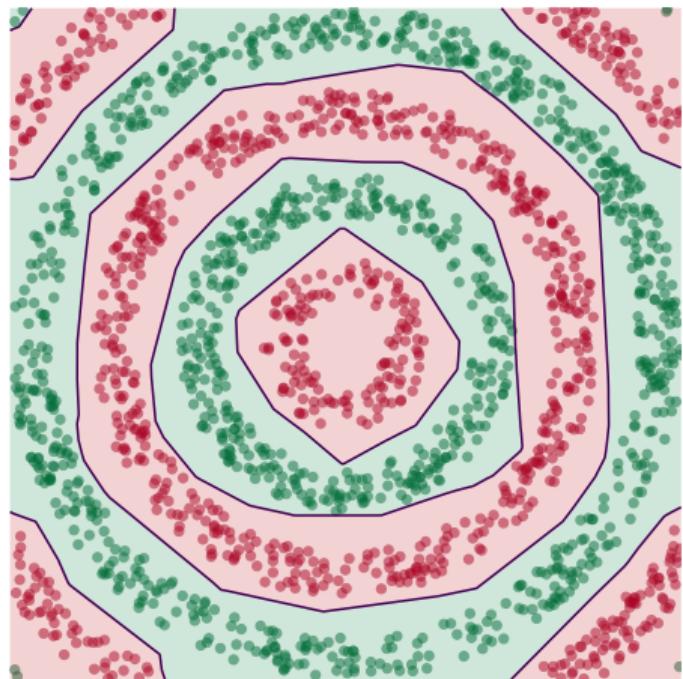
Learning a New Representation



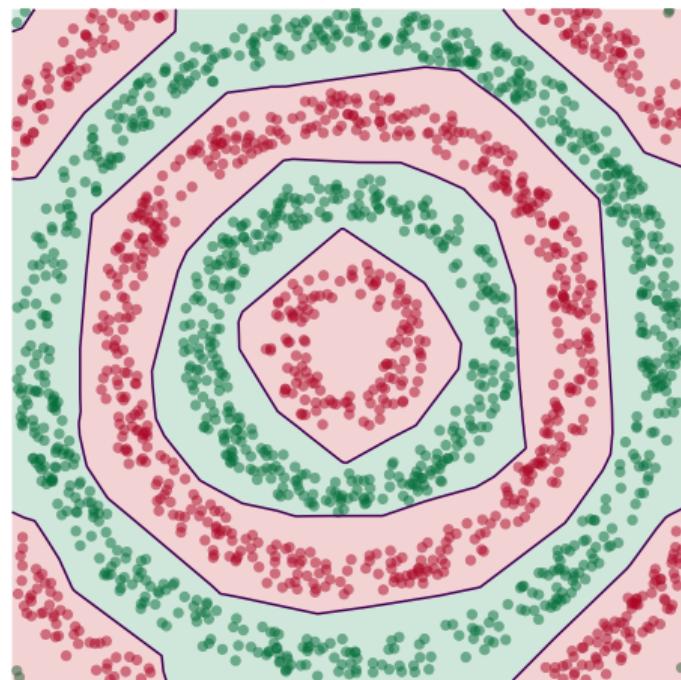
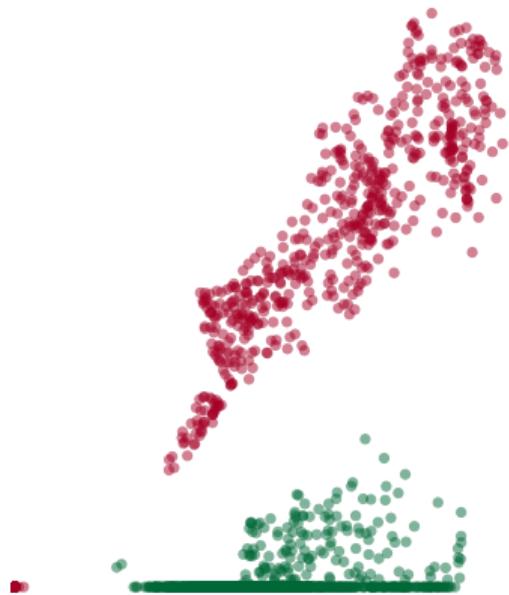
Learning a New Representation



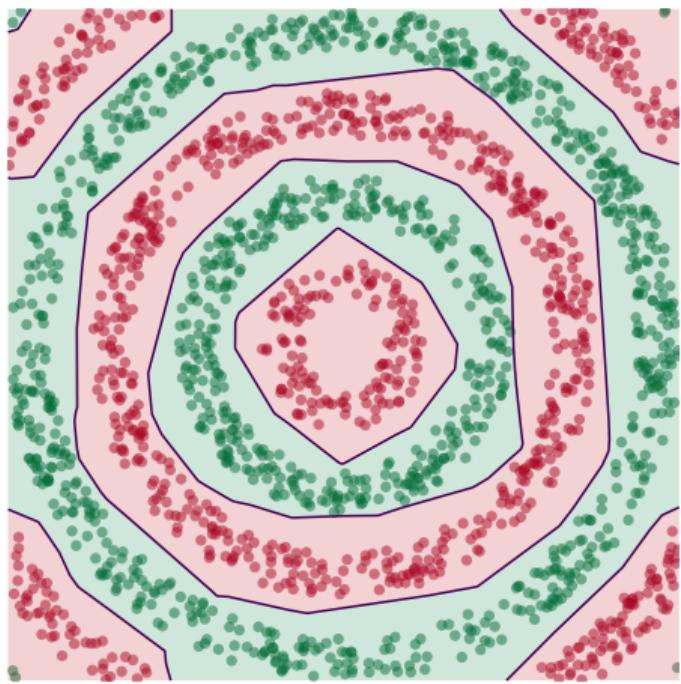
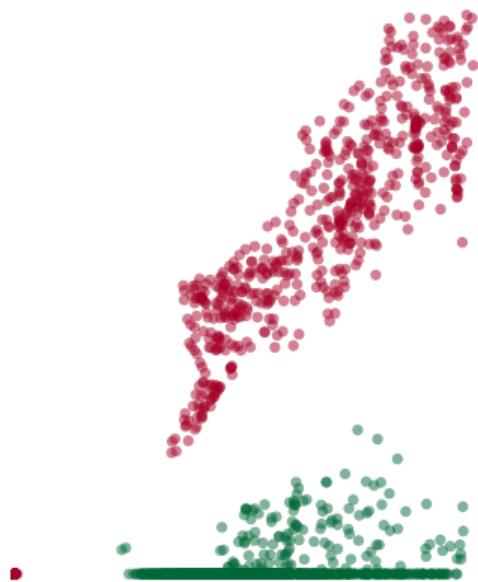
Learning a New Representation



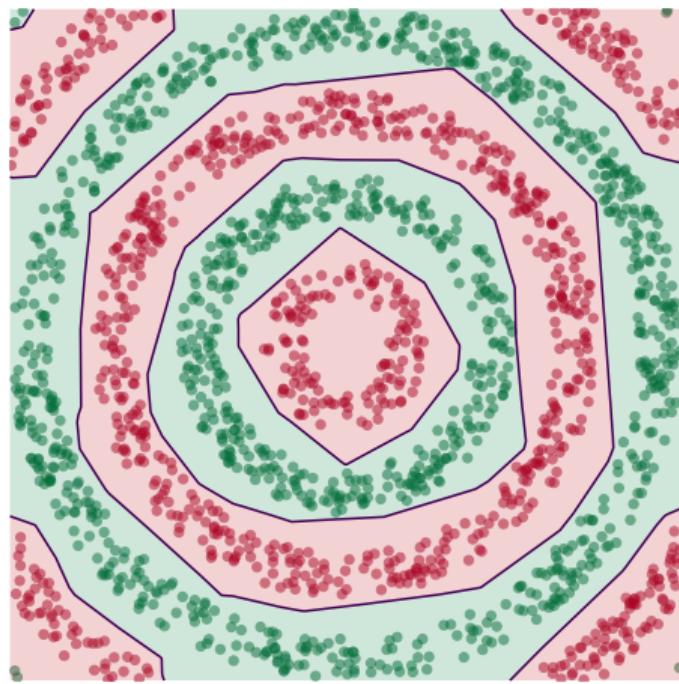
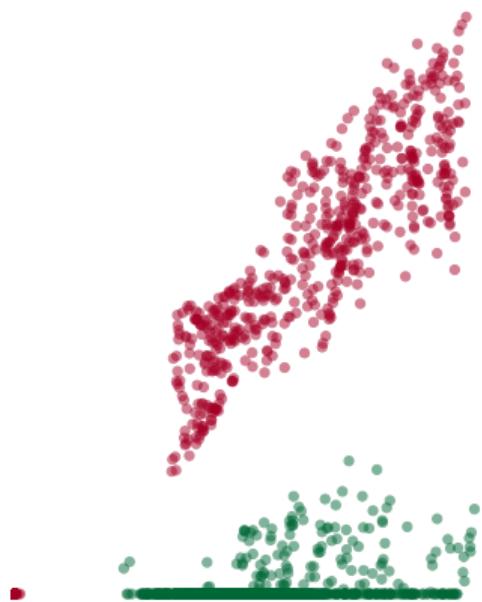
Learning a New Representation



Learning a New Representation



Learning a New Representation



Deep Networks and Approximation

- ▶ Deep networks are also universal approximators.
- ▶ May require fewer nodes and/or parameters than single hidden layer.
- ▶ I.e., there exist functions which require an exponential number of nodes to approximate with a single hidden layer, but not with several layers.

Challenges

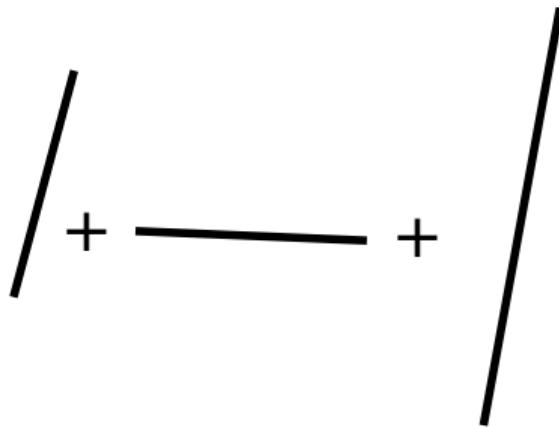
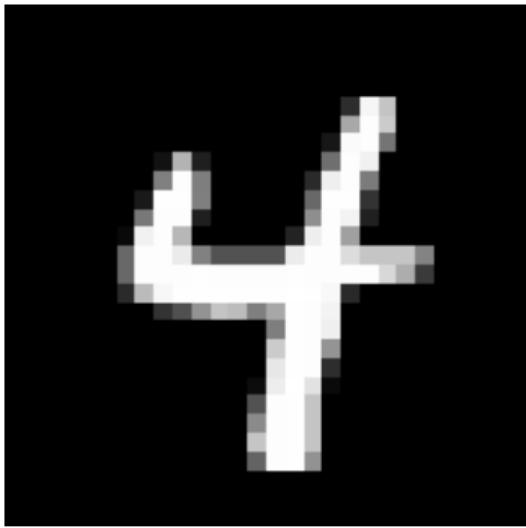
- ▶ The deeper the network, the weaker the gradient gets.
- ▶ Very non-convex!
- ▶ Deeper networks are harder to learn.

DSC 190

Machine Learning: Representations

Lecture 15 | Part 4

Convolutions

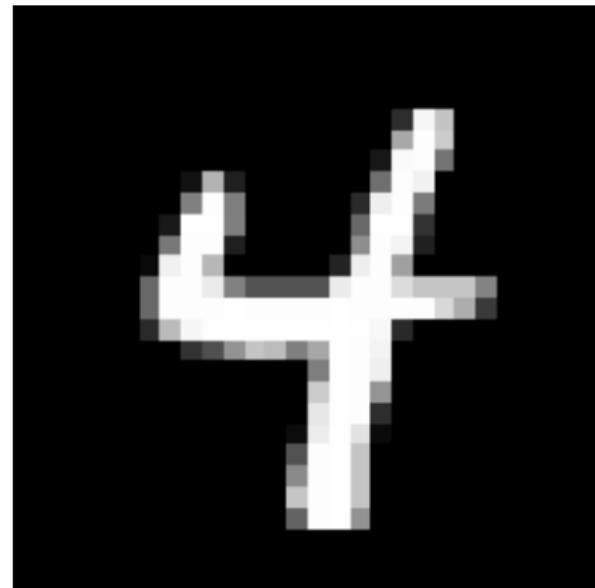


From Simple to Complex

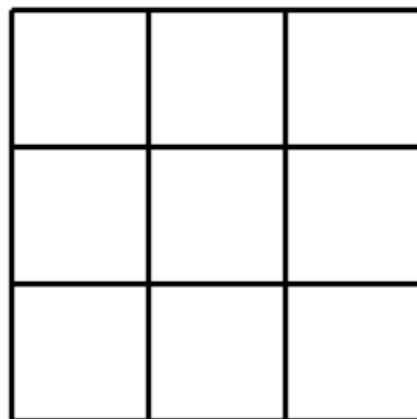
- ▶ Complex shapes are made of simple patterns
- ▶ The human visual system uses this fact
- ▶ Line detector → shape detector → ... → face detector
- ▶ Can we replicate this with a deep NN?

Edge Detector

- ▶ How do we find **vertical edges** in an image?
- ▶ One solution: **convolution** with an **edge filter**.



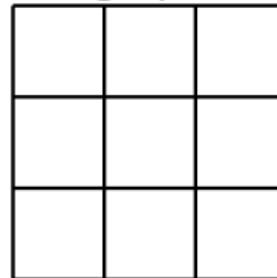
Vertical Edge Filter



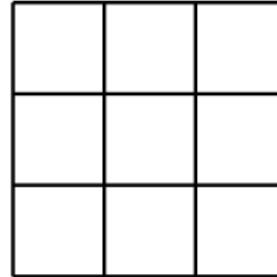
Idea

- ▶ Take a patch of the image, same size as filter.
- ▶ Perform “dot product” between patch and filter.
- ▶ If large, this is a (vertical) edge.

image patch:



filter:



Idea

- ▶ Move the filter over the entire image, repeat procedure.

$$\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 0 & .9 \\ \hline 0 & 0 & .9 \\ \hline 0 & 0 & .8 \\ \hline 0 & 0 & .7 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline \end{array} = \begin{array}{|c|c|c|} \hline \end{array}$$

The diagram illustrates the convolution process. On the left is a 5x3 input matrix with values 0, 0, 0; 0, 0, 0.9; 0, 0, 0.9; 0, 0, 0.8; and 0, 0, 0.7. A 3x3 filter matrix, consisting of three rows and three columns of empty boxes, is positioned to the right of the input. An asterisk (*) symbol indicates the multiplication operation, and an equals sign (=) indicates the result. To the right of the equals sign is a 5x5 output matrix, also consisting of empty boxes, representing the result of the convolution step.

Idea

- ▶ Move the filter over the entire image, repeat procedure.

$$\begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 & \dots \\ 0 & 0 & .9 & 0 & 0 & 0 & \dots \\ 0 & 0 & .9 & 0 & 0 & 0 & \dots \\ 0 & 0 & .8 & 0 & 0 & 0 & \dots \\ 0 & 0 & .7 & 0 & 0 & 0 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots \end{matrix} * \begin{matrix} \boxed{} & \boxed{} & \boxed{} \\ \boxed{} & \boxed{} & \boxed{} \\ \boxed{} & \boxed{} & \boxed{} \end{matrix} = \begin{matrix} \boxed{} & \boxed{} & \boxed{} & \boxed{} & \boxed{} & \dots \\ \boxed{} & \boxed{} & \boxed{} & \boxed{} & \boxed{} & \dots \\ \boxed{} & \boxed{} & \boxed{} & \boxed{} & \boxed{} & \dots \\ \boxed{} & \boxed{} & \boxed{} & \boxed{} & \boxed{} & \dots \\ \boxed{} & \boxed{} & \boxed{} & \boxed{} & \boxed{} & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots \end{matrix}$$

Idea

- ▶ Move the filter over the entire image, repeat procedure.

$$\begin{array}{|c|c|c|c|c|c|} \hline 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & .9 & 0 & 0 & .7 \\ \hline 0 & 0 & .9 & 0 & 0 & .8 \\ \hline 0 & 0 & .8 & 0 & 0 & .9 \\ \hline 0 & 0 & .7 & 0 & 0 & 0 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline \end{array} = \begin{array}{|c|c|c|c|c|c|} \hline \end{array}$$

Idea

- ▶ Move the filter over the entire image, repeat procedure.

$$\begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & .9 & 0 & 0 & .7 \\ 0 & 0 & .9 & 0 & 0 & .8 \\ 0 & 0 & .8 & 0 & 0 & .9 \\ 0 & 0 & .7 & 0 & 0 & 0 \end{matrix} * \begin{matrix} \quad & \quad & \quad \\ \quad & \quad & \quad \\ \quad & \quad & \quad \end{matrix} = \begin{matrix} \quad & \quad & \quad & \quad & \quad & \quad \\ \quad & \quad & \quad & \quad & \quad & \quad \\ \quad & \quad & \quad & \quad & \quad & \quad \\ \quad & \quad & \quad & \quad & \quad & \quad \\ \quad & \quad & \quad & \quad & \quad & \quad \\ \quad & \quad & \quad & \quad & \quad & \quad \end{matrix}$$

Idea

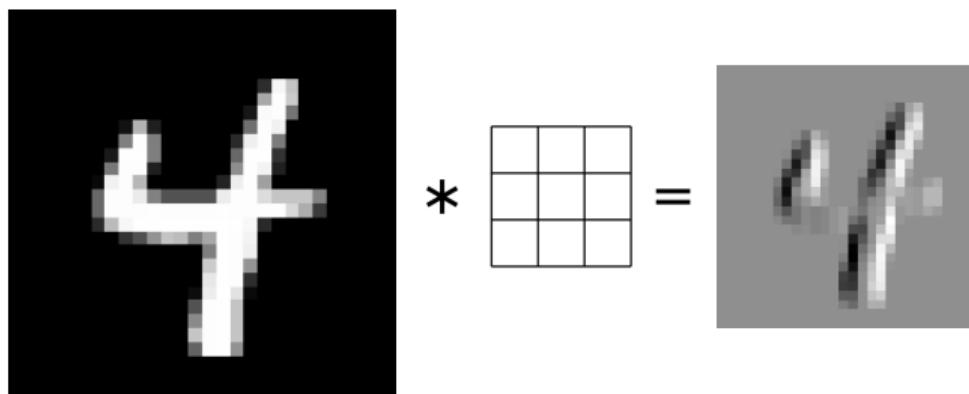
- ▶ Move the filter over the entire image, repeat procedure.

$$\begin{array}{|c|c|c|c|c|c|c|} \hline 0 & 0 & 0 & 0 & 0 & 0 & \\ \hline 0 & 0 & .9 & 0 & 0 & 0 & .7 \\ \hline 0 & 0 & .9 & 0 & 0 & 0 & .8 \\ \hline 0 & 0 & .8 & 0 & 0 & 0 & .9 \\ \hline 0 & 0 & .7 & 0 & 0 & 0 & 0 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline \end{array} = \begin{array}{|c|c|c|c|c|c|c|} \hline \end{array}$$

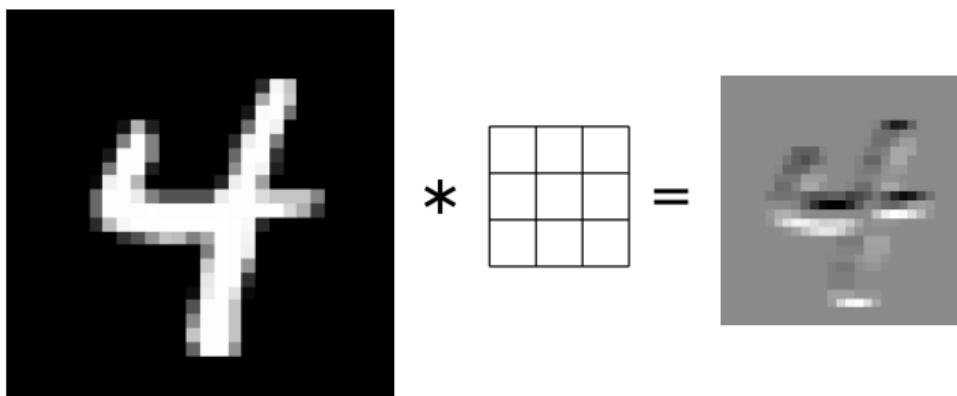
Convolution

- ▶ The result is the (2d) **convolution** of the filter with the image.
- ▶ Output is also 2-dimensional array.
- ▶ Called a **response map**.

Example: Vertical Filter



Example: Horizontal Filter



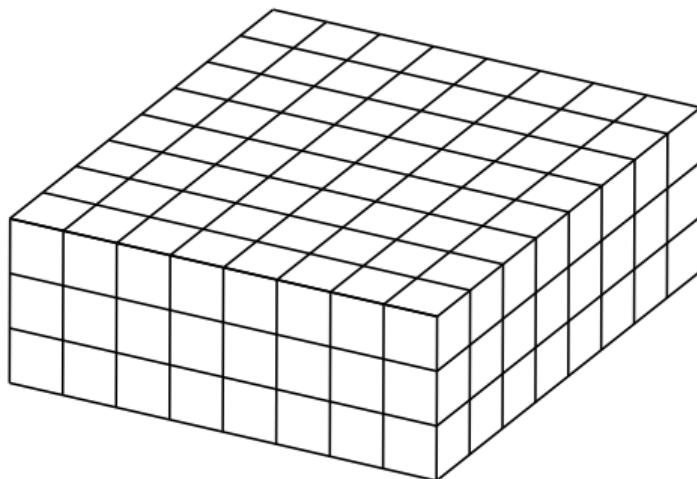
More About Filters

- ▶ Typically 3×3 or 5×5 .
- ▶ Variations: different **stride**, image **padding**.

3-d Filters

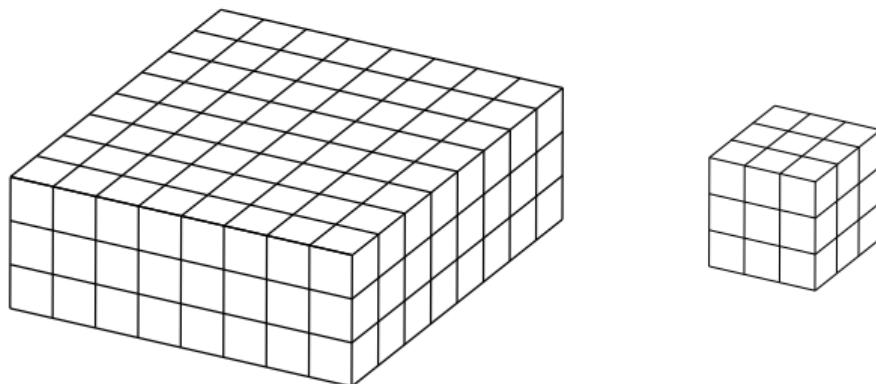
- ▶ Black and white images are 2-d arrays.
- ▶ But color images are 3-d arrays:
 - ▶ a.k.a., **tensors**
 - ▶ Three color **channels**: red, green, blue.
 - ▶ $\text{height} \times \text{width} \times 3$
- ▶ How does convolution work here?

Color Image

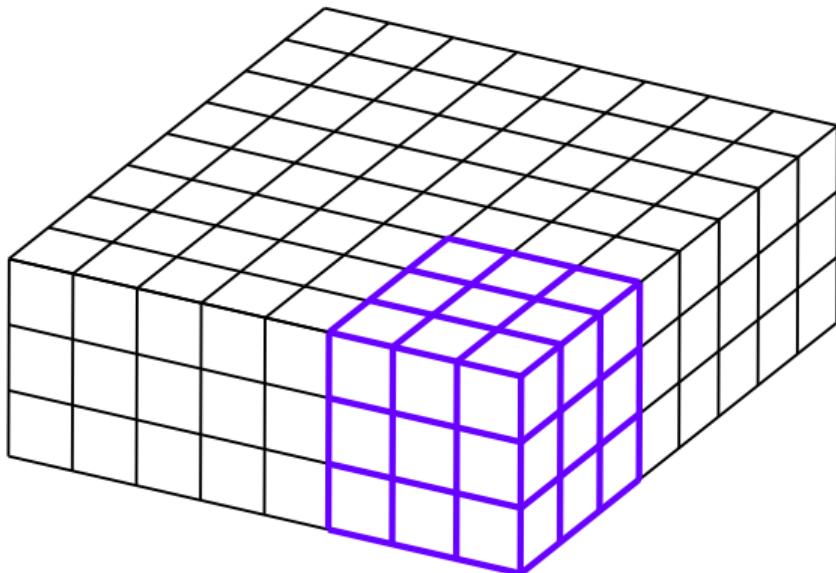


3-d Filter

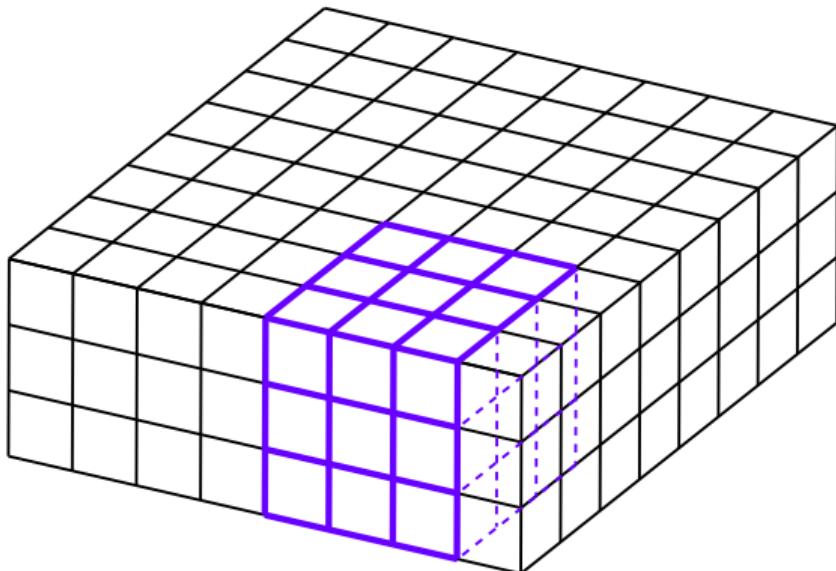
- ▶ The filter must also have three channels:
 - ▶ $3 \times 3 \times 3$, $5 \times 5 \times 3$, etc.



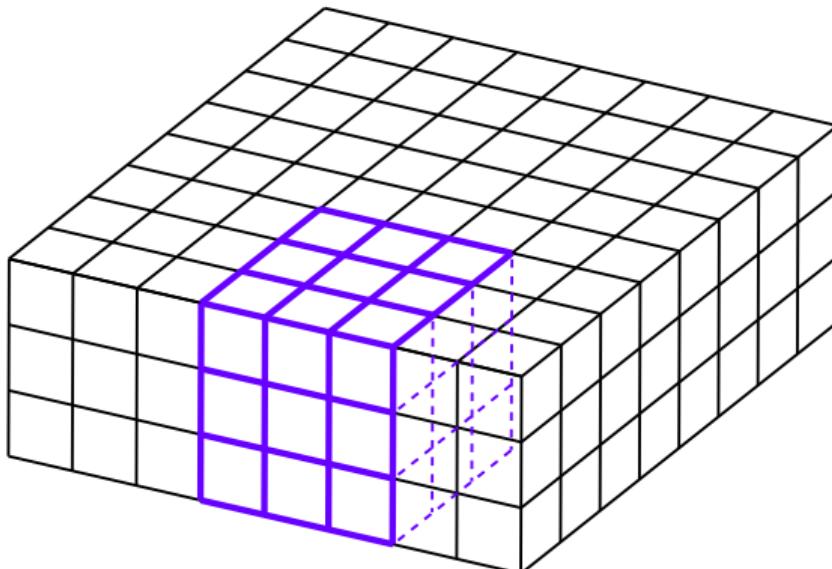
3-d Filter



3-d Filter



3-d Filter



Convolution with 3-d Filter

- ▶ Filter must have same number of channels as image.
 - ▶ 3 channels if image RGB.
- ▶ Result is still a 2-d array.

General Case

- ▶ Input “image” has k channels.
- ▶ Filter must have k channels as well.
 - ▶ e.g., $3 \times 3 \times k$
- ▶ Output is still $2 - d$

